



# SOCKET PROGRAMMING IN C/C++




# SOCKETS

## Definition & Characteristics




# Introduction

Sockets are a protocol independent method of creating a connection between processes. Sockets can be either



- connection based or connectionless:  
Is a connection established before communication or does each packet describe the destination?
- packet based or streams based:  
Are there message boundaries or is it one stream?
- reliable or unreliable.  
Can messages be lost, reordered, or corrupted?





# Socket characteristics


Domain, type and transport protocol.

Common domains

- **AF\_INET** (ARPA Internet protocols)
- **AF\_UNIX** (UNIX internal protocols)
- **AF\_ISO** (ISO protocols)
- **AF\_NS** (Xerox Network Systems protocols)

Common types

- **SOCK\_STREAM** (virtual circuit: received in order, reliably)
- **SOCK\_DGRAM** (datagram: arbitrary order, unreliable)




# Socket characteristics (cont'd)

Each socket type has one or more protocols. Ex:

- **TCP/IP** (virtual circuits)
- **UDP** (datagram)

Use of sockets:

- Connection-based sockets communicate client-server: the server waits for a connection from the client
- Connectionless sockets are peer-to-peer: each process is symmetric.



# Socket Functions

socket: creates a socket of a given domain, type, protocol (buy a phone)

bind: assigns a name and an address to the socket (get a telephone number)

listen: specifies the number of pending connections that can be queued for a server socket. (call waiting allowance)

accept: server accepts a connection request from a client (answer phone)

connect: client requests a connection request to a server (call)


send, sendto: write to connection (speak)

recv, recvfrom: read from connection (listen)

close: end the call







# Socket Functions

socket: creates a socket of a given domain, type, protocol (buy a phone)

bind: assigns a name and an address to the socket (get a telephone number)

listen: specifies the number of pending connections that can be queued for a server socket. (call waiting allowance)

accept: server accepts a connection request from a client (answer phone)


connect: client requests a connection request to a server (call)

send, sendto: write to connection (speak)


recv, recvfrom: read from connection (listen)

close: end the call






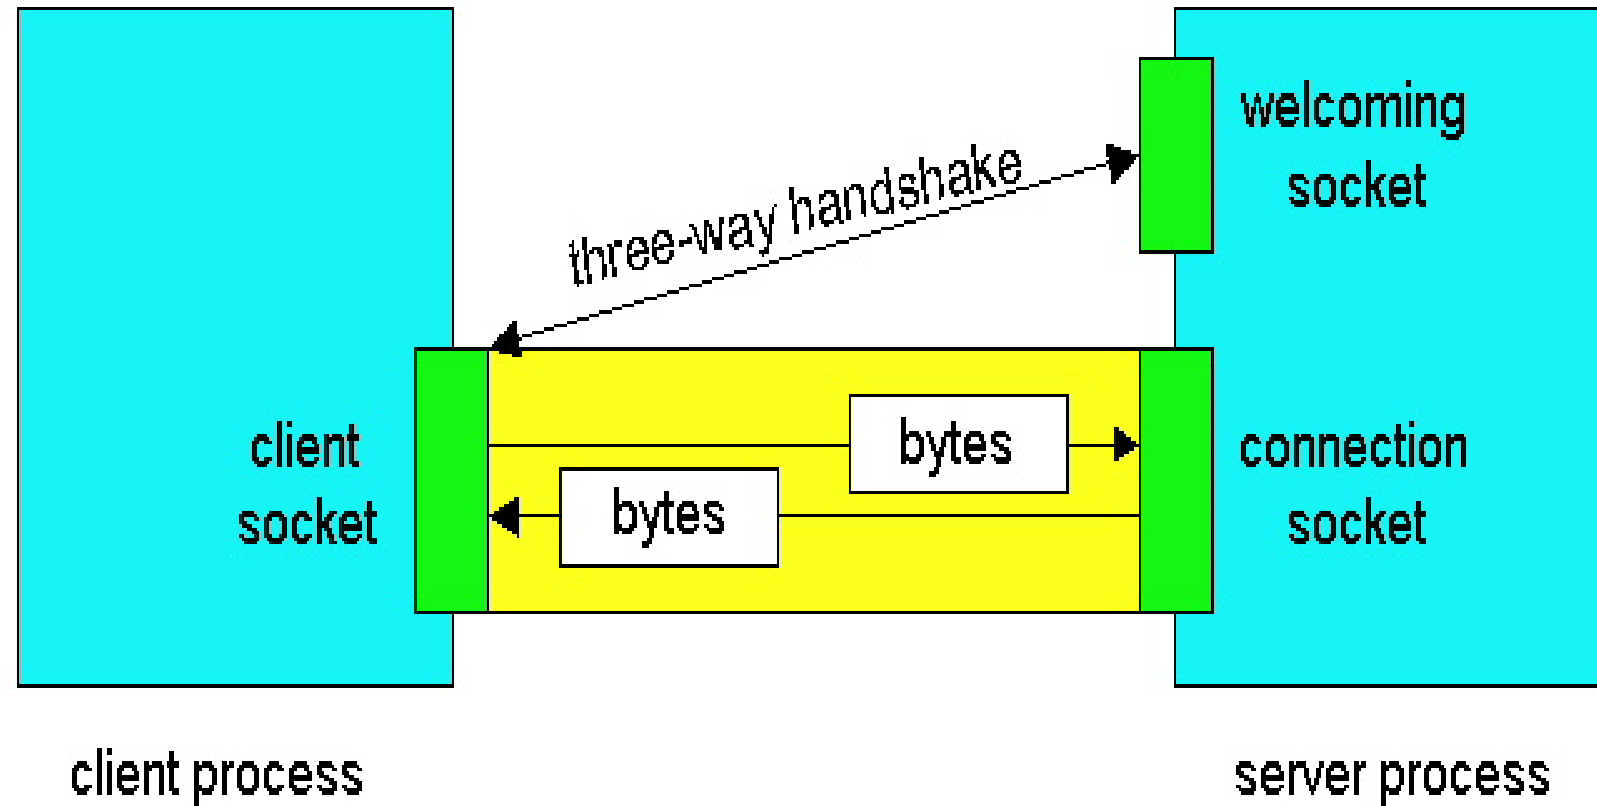
# TCP-based sockets










# Accept (three-way handshake)





# UDP-based sockets






# SOCKET API

## Functions






# socket

**`int socket ( int domain , int type , int protocol ) ;`**

Returns a file descriptor (socket ID) if successful, -1 otherwise.

- domain is AF\_INET for internet communication to IP addresses
- The type argument can be:
  - SOCK\_STREAM: Establishes a virtual circuit for stream
  - SOCK\_DGRAM: Establishes a datagram for communication
  - SOCK\_RAW (IP level).
- protocol specifies the protocol used. It is usually 0 to use the default protocol for the chosen domain and type.




# bind

**int bind ( int sid , struct sockaddr \*addrPtr , int len )**

- sid: is the socket id
- addrPtr: is a pointer to the address family dependent address structure
- len: is the size of \*addrPtr

Associates a socket id with an address to which other processes can connect. In internet protocol the address is [ipNumber, portNumber]




# connect

*this is the first of the client calls*

**int connect ( int sid, struct sockaddr \* addrPtr, int len )**

- Specifies the destination to form a connection with (addrPtr), and returns a 0 if successful, -1 otherwise.






# Port usage

Note that the initiator of communications needs a fixed port to target communications.

This means that some ports must be reserved for these “well known” ports.

Port usage:


- 0-1023: These ports can only be bound by root
- 1024-49151: well known ports
- 49152-65535: ephemeral (short-term) ports for OS



# listen

**int listen ( int sid , int size )**

- Where size is the number of pending connection requests allowed (typically limited by Unix kernels to 5).
- Returns 0 on success, -1 on failure.




# accept

**int accept ( int sid, struct sockaddr \*addrPtr, int\* lenPtr )**

- Returns the socketId and address of client connecting to socket. If lenPtr or addrPtr equal zero, no address structure is returned. lenPtr is the maximum size of address structure that can be called, returns the actual value.
- Waits for an incoming request, and when received creates a socket for it.






# send

**int send ( int sid, const char \* bufferPtr, int len, int flag )**


- Send a message. Returns the number of bytes sent or -1 if failure. (Must be a bound socket).



# sendto

**int sendto ( int sid, const void \*bufferPtr, size\_t bufferLength,  
int flag, struct sockaddr \*addrPtr, socklen\_t addrLength )**

- for connectionless protocols
- sends a buffer, at bufferPtr, of length bufferLength to the address specified by addrPtr of size addrLength.
- returns number of bytes sent or -1 on error.




# recv

**int recv ( int sid, const char \* bufferPtr, int len, int flag )**

- Receive up to len bytes in bufferPtr. Returns the number of bytes received or -1 on failure.






# recvfrom

```
int recvfrom ( int sid, void *bufferPtr, int bufferLength, int flag,  
sockaddr *addrPtr, int *addrLengthPtr )
```


- for connectionless protocols
- receive a buffer in bufferPtr of maximum length bufferLength from an unspecified sender.
- sender address returned in addrPtr, of size \*addrLengthPtr.
- returns number of bytes received or -1 on error.



# close

```
int close ( int sid );
```


Both client and server close the sockets after data transfer



# PROGRAMMING

APIs and Examples






# APIs for managing names and IP addresses

A number of auxiliary APIs used:

- The hostent structure: describes IP, hostname pairs
- gethostbyname: hostent of a specified machine
- htons, htonl, ntohs, ntohl: byte ordering
- inet\_pton, inet\_ntop, inet\_addr: conversion of IP numbers between presentation and strings




# sockaddr

For the internet family:

```
struct in_addr {  
    long s_addr;  
};
```

```
struct sockaddr_in {  
    int sin_family;           // AF_INET identifiers  
    int sin_port;             // port number,  
                                // if 0 then kernel chosen  
    struct in_addr sin_addr; // IP address  
                                // INADDR_ANY refers to the IP  
                                // addresses of the current host  
    char sin_zero[8];         // Unused, always zero  
};
```

*server.sin\_addr.s\_addr=inet\_addr("127.0.0.1");*




# hostent

```
struct hostent {  
    char*  h_name;    //official (canonical) name of the host  
    char **h_aliases; //null terminated array of alternative hostnames  
    int    h_addrtype; //host address type AF_INET or AF_INET6  
    int    h_length;   //4 or 16 bytes  
    char **h_addr_list; //IPv4 or IPv6 list of address  
}
```

- Error is return through h\_error which can be:
  - HOST NOT FOUND
  - TRY AGAIN
  - NO RECOVERY
  - NO DATA





# gethostbyname


**Struct hostent\* gethostbyname (const char\* hostname)**

- Auxiliary functions
- Translates a DNS name into a hostent.

- Example:

```
Struct hostent* hostEntity = gethostbyname("eee.metu.edu.tr");
```

```
Memcpy (socketAddr->sin_addr, hostEntity->h_addr_list[0], hostEntity->h_length);
```



# Network byte ordering

`unsigned long htonl(unsigned long n)`

`// host to network conversion of a 32-bit value;`

`unsigned short htons(unsigned short n)`


`// host to network conversion of a 16-bit value;`

`unsigned long ntohl(unsigned long n)`

`// network to host conversion of a 32-bit value;`

`unsigned short ntohs(unsigned short n)`

`// network to host conversion of a 16-bit value.`



# IP Number translation


`inet_pton`, `inet_ntop`, `inet_addr` routines translate between the address as a string and the address as the number.

We have 4 representations:

- IP number in host order
- IP number in network order
- Presentation (ex. dotted decimal)
- Fully qualified domain name

Only the last needs an outside lookup to convert to one of the other formats.






# Example: TCP/IP Server Code

Without error checking.


```
sockaddr_in serverAddr ;
sockaddr &serverAddrCast = (sockaddr &) serverAddr ;
    // get a tcp/ip socket
serverAddr.sin_family = AF_INET;
    // get an address from any internet interface on this server.
serverAddr.sin_addr.s_addr = htonl (INADDR_ANY) ;
serverAddr.sin_port = htons(13) ;
int listenFd = socket (AF_INET , SOCK_STREAM, 0 ) ;
bind (listenFd , &serverAddrCast, sizeof ( serverAddr ) ) ;
Listen (listenFd, 5 );
for ( ; ; ){
int connectFd = accept (listenFd, (sockaddr*) NULL, NULL ) ;
    // read and write operations on connectFd
close(connectFd) ;
}
```



# Example: TCP/IP Client Code

Without error checking.

```
sockaddr_in serverAddr ;
sockaddr &serverAddrCast = (sockaddr &) serverAddr ;
    // get a tcp/ip socket
serverAddr.sin_family = AF_INET;
    // host IP # in dotted decimal format !
serverAddr.sin_addr.s_addr = inet_addr("144.122.167.37");
serverAddr.sin_port = htons(13) ;
connect (sockFd, serverAddrCast, sizeof (serverAddr));
    // .. read and write operations on sockFd ..
close(connectFd ) ;
}
```




# Example: UDP Server

```
int socketId = socket (AF_INET , SOCK_DGRAM, 0 ) ;
sockaddr_in serverAddr , clientAddr ;
sockaddr &serverAddrCast = ( sockaddr &) serverAddr ;
sockaddr &clientAddrCast = ( sockaddr &) clientAddr ;
// allow connection to any addr on host for hosts with multiple network connections and ast server port .
serverAddr.sin_family = AF_INET ;
serverAddr.sin_port = htons ( serverPort ) ;
serverAddr.sin_addr.s_addr = INADDR_ANY;
bind ( socketId, &serverAddrCast, sizeof (addr) ) ;


int size = sizeof(clientAddr);
recvfrom ( socketId , buffer , bufferSize , 0 , clientAddrCast , &size ) ;
sendto ( socketId , buffer , bufferSize , 0 , clientAddrCast , size ) ; // reply to the client just received from
close(socketId);
```






# Example: UDP Client

```
int socketId = socket (AF_INET , SOCK_DGRAM, 0 ) ;
sockaddr_in serverAddr , clientAddr ;
sockaddr &serverAddrCast = ( sockaddr &) serverAddr ;
sockaddr &clientAddrCast = ( sockaddr &) clientAddr ;
// specify server address and port
serverAddr.sin_family = AF_INET ;
serverAddr.sin_port = htons ( serverPort ) ;
struct hostent *hp = gethostbyname ( hostName ) ;
memcpy ((char *) &serverAddr.sin_addr, ( char *)hp->h_addr , hp->h_length ) ;
// no need to bind if not peer-to-peer
int size = sizeof(serverAddr);
sendto ( socketId , buffer , bufferSize , 0 , serverAddrCast , size ) ;
Recvfrom (socketId , buffer , bufferSize , 0 , serverAddrCast , &size );
close(socketId);
```



# Timeout for Blocking Functions

- Recv() or recvfrom() waits until some data arrives
- If arriving packet is lost (especially for UDP), recvfrom() function waits forever
- This can be avoided with timeout functions
  - Select, poll, alarm are some example functions.



# THANKS

Adopted from; **Mani Radhakrishnan and Jon Solworth**  
**Socket Programming in C/C++**; September 24, 2004

Prepared by **Yiğit Özcan**; April, 2013

Edited by **Utku Civelek**; April, 2016