

# **Analytical Modeling of Parallel Systems and Performance Analysis**



# Topic Overview

- Sources of Overhead in Parallel Programs
- Performance Metrics for Parallel Systems
- Effect of Granularity on Performance
- Scalability of Parallel Systems

# Analytical Modeling - Basics

- A sequential algorithm is evaluated by its runtime
  - In general, asymptotic runtime as a function of input size
- The asymptotic runtime is independent of the platform.
  - Analysis “at a constant factor”.
- A parallel algorithm has more parameters.

# Big O notation – $O(g(n))$

- Big O notation is used to describe the performance or complexity of an algorithm.
- Big O specifically describes the bounds
  - can be thought of as **worst-case** scenario
  - can be used to describe the execution time required by an algorithm.

# O(1)

- O(1) describes an algorithm that will always execute in the same time regardless of the size of the input data set.

```
bool IsFirstElementNull(String[] strings)
{
    if(strings[0] == null)
        return true;
    return false;
}
```

# $O(n)$

- $O(n)$  describes an algorithm whose performance will grow linearly and in direct proportion to the size of the input data set.

```
bool ContainsValue(String[] strings, String value)
{
    for(int i = 0; i < strings.Length; i++)
    {
        if(strings[i] == value)
            return true;
    }
    return false;
}
```


# $O(n^2)$

- $O(n^2)$  represents an algorithm whose performance is directly proportional to the square of the size of the input data set.
- This is common with algorithms that involve nested iterations over the data set. Deeper nested iterations will result in  $O(n^3)$ ,  $O(n^4)$  etc.

```
bool ContainsDuplicates(String[] strings)
{
    for(int i = 0; i < strings.Length; i++)
    {
        for(int j = 0; j < strings.Length; j++)
        {
            if(i == j) // Don't compare with self
                continue;
            if(strings[i] == strings[j])
                return true;
        }
    }
    return false;
}
```

# $O(\log n)$


- Algorithms iteratively halving the datasets as we have seen before.
- Example: Parallel Reduction





Execution  
time

$$O(g(n))$$



# Big-Theta ( $\Theta$ ) and Big-Omega ( $\Omega$ )

- $f(n) \in O(g(n))$ :
  - *f is bounded above by g asymptotically*
  - (worst case scenario)
- $f(n) \in \Omega(g(n))$  :
  - *f is bounded below by g asymptotically*
  - (best case scenario)
- $f(n) \in \Theta(g(n))$ :
  - *f is bounded both above and below*

# Analytical Modeling - Basics

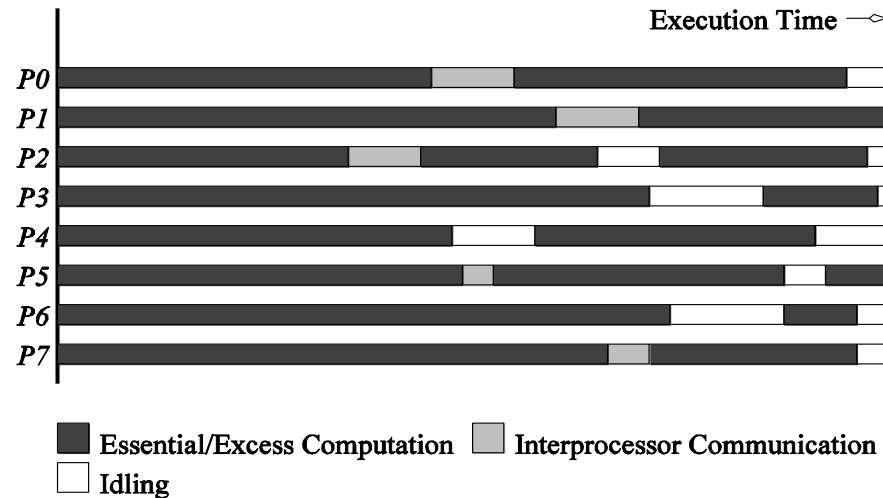
- The parallel runtime of a program depends on the:
  - input size,
  - the number of processors,
  - the communication parameters of the machine.
- An algorithm must therefore be analyzed in the context of the underlying platform.
- A parallel system is a combination of:
  - a parallel algorithm
  - an underlying platform.

# Analytical Modeling - Basics

- A number of performance measures are intuitive.
- Execution time: the time from the start of the **first** processor to the stopping time of the **last** processor in a parallel ensemble ( $T_p$ ).
- But how does this scale when the number of processors is changed and the program is ported to another machine altogether?
- How much faster is the parallel version?
- This brings the obvious follow up question:  
“What’s the baseline serial version with which we compare?”

# Sources of Overhead in Parallel Programs

- If I use two processors, shouldn't my program run twice as fast?  
No - a number of overheads, including wasted computation, communication, idling, and contention cause degradation in performance.



The execution profile of a hypothetical parallel program executing on eight processing elements. Profile indicates times spent performing computation (both essential and excess), communication, and idling.

# Sources of Overheads in Parallel Programs

- **Inter-process interactions:** Processors working on any non-trivial parallel problem will need to talk to each other.
- **Idling:** Processes may idle because of load imbalance, synchronization, or serial components.
- **Excess Computation:** This is computation not performed by the serial version.
  - This might be because the serial algorithm is difficult to parallelize,
  - or that some computations are repeated across processors to minimize communication.

# Performance Metrics for Parallel Systems:

## Execution Time

- Serial runtime of a program is the time elapsed between the beginning and the end of its execution on a sequential computer.
- The parallel runtime is the time that elapses from the moment the first processor starts to the moment the last processor finishes execution.
- We denote the serial runtime by  $T_s$  and the parallel runtime by  $T_p$ .

# Performance Metrics for Parallel Systems:

## Total Parallel Overhead

- Let  $T_{all}$  be the total time collectively spent by all the processing elements.

$$T_{all} = p T_p \quad (p \text{ is the number of processors}).$$

- Observe that  $T_{all} - T_s$  is the total time spend by all processors combined in non-useful work. This is called the *total overhead* ( $T_o$ ).

$$T_o = p T_p - T_s$$



# Performance Metrics for Parallel Systems:


## Speedup

- What is the benefit from parallelism?
- Speedup ( $S$ ) is the ratio of the time taken to solve a problem on a single processor to the time required to solve the same problem on a parallel computer with  $p$  identical processing elements.


## Performance Metrics: Example

- Consider the problem of adding  $n$  numbers by using  $n/2$  processing elements.
- If  $n$  is a power of two, we can perform this operation in  $\log_2 n$  steps by parallel reduction


# Performance Metrics: Example




(a) Initial data distribution and the first communication step




(b) Second communication step



(c) Third communication step



(d) Fourth communication step



(e) Accumulation of the sum at processing element 0 after the final communication

Computing the global sum of 16 partial sums using 8 processing elements .

# Performance Metrics: Example

- We have the parallel time

$$T_p = \Theta(\log_2 n)$$

- We know that  $T_s = \Theta(n)$

- Speedup  $S$  is given by  $S = \Theta(n / \log_2 n)$

# Performance Metrics: Speedup

- For a given problem, there might be many serial algorithms available.
- These algorithms may have different asymptotic runtimes and may be parallelizable to different degrees.
- For the purpose of computing speedup, we always consider **the best sequential program** as the baseline.

# Performance Metrics: Speedup Example

Consider the problem of parallel bubble sort.

- The serial time for bubble sort is 150 seconds.
- The parallel time for odd-even sort (efficient parallelization of bubble sort) is 40 seconds.
  - The speedup would appear to be  $150/40 = 3.75$ .
- What if another serial quicksort implementation only took 30 seconds?
  - In this case, the speedup is  $30/40 = 0.75$ .

# Performance Metrics: Speedup Bounds

- Speedup can be as low as 0 (the parallel program never terminates).
- Speedup, in theory, should be upper bounded by  $p$ 
  - we can only expect a  $p$ -fold speedup if we use  $p$  processing elements.

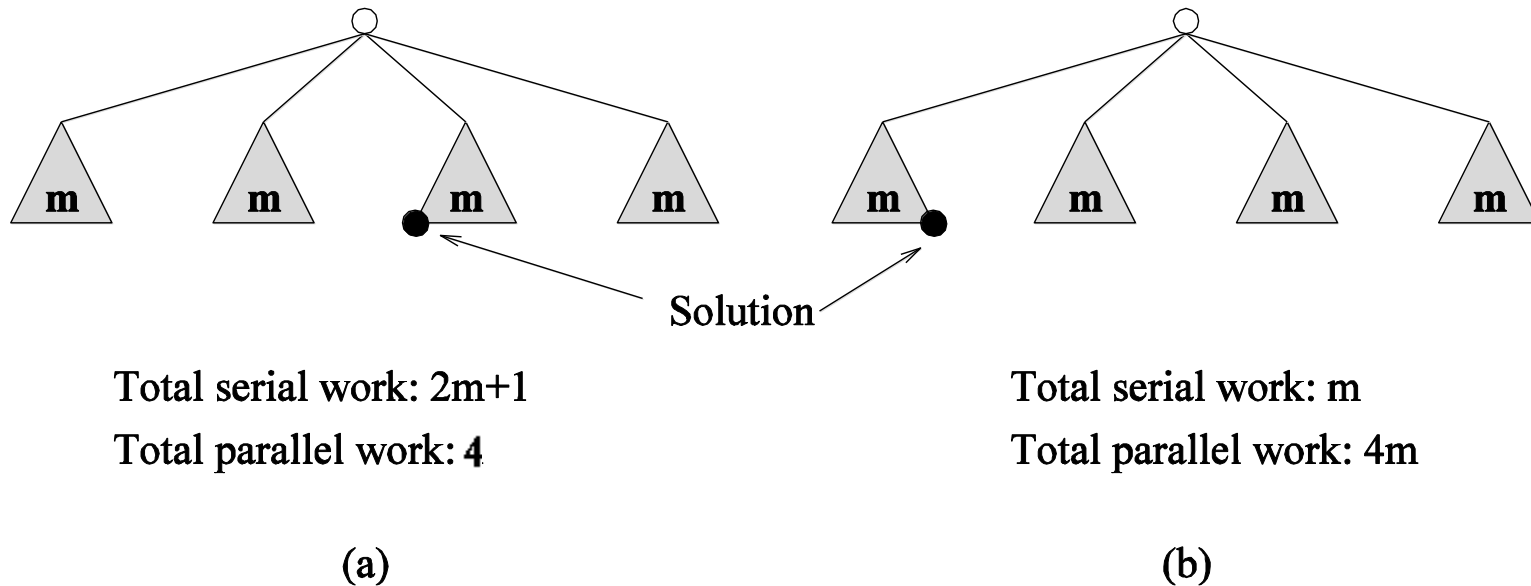
# Performance Metrics: Speedup Bounds

- A speedup greater than  $p$  is possible only if each processing element spends less than time  $T_s/p$  solving the problem.
  - In this case, a single processor could be time-slided to achieve a faster serial program, which contradicts our assumption of fastest serial program as basis for speedup.



# Performance Metrics: Superlinear Speedups

One reason for superlinearity is that the parallel version does less work than corresponding serial algorithm.



(a) Shows super-linear behavior while (b) shows sub-linear behavior

# Performance Metrics: Superlinear Speedups

Resource-based superlinearity:

The higher aggregate cache/memory bandwidth can result in better cache-hit ratios, and therefore superlinearity.

Example:

- A processor with 64KB of cache yields an 80% hit ratio, the remaining comes from local memory.
- If two processors are used, since the problem size/processor is smaller, the hit ratio goes up to 90%. Of the remaining 10% access, 8% come from local memory and 2% from remote memory.

If DRAM access time is 100 ns, cache access time is 2 ns, and remote memory access time is 400ns,

Case1:  $2 \times 0.8 + 100 \times 0.2 = 21.6$  ns

Case 2:  $2 \times 0.9 + 100 \times 0.08 + 400 \times 0.02 = 17.8$  ns

This corresponds to a speedup of 1.21 in memory access.

# Performance Metrics: Superlinear Speedups

Example:

DRAM access time is 100 ns

cache access time is 2 ns

remote memory access time is 400ns

Case1: A processor with 64KB of cache yields an 80% hit ratio, the remaining comes from local memory.

$$2*0.8+100*0.2=21.6 \text{ ns}$$

Case 2: Two processors are used, since the problem size/processor is smaller, the hit ratio goes up to 90%. Of the remaining 10% access, 8% come from local memory and 2% from remote memory.

$$2*0.9+100*0.08+400*0.02=17.8 \text{ ns}$$

This corresponds to a speedup of 1.21 in memory access.

# Performance Metrics: Efficiency

- Efficiency is a measure of the fraction of time for which a processing element is usefully employed
- Mathematically, it is given by

$$E = \frac{S}{p}, \quad 0 \leq E \leq 1$$

# Performance Metrics: Efficiency

## Example

- The speedup of adding numbers on  $n$  processors is given by

$$S = \frac{n}{\log n}$$

- Efficiency is given by

$$\begin{aligned} E &= \frac{\Theta\left(\frac{n}{\log n}\right)}{n} \\ &= \Theta\left(\frac{1}{\log n}\right) \end{aligned}$$

# Parallel Time, Speedup, and Efficiency Example

Consider the problem of filtering images.

The problem requires us to apply a template to each pixel.

random image  $I(x,y)$

8	8	2	2	12
1	3	4	7	7
3	15	5	9	5
3	1	9	12	12
1	3	15	4	15

averaging filter  
 $W(x,y)$

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

$$\frac{8}{9} + \frac{8}{9} + \frac{2}{9} + \frac{1}{9} + \frac{3}{9} + \frac{4}{9} + \frac{3}{9} + \frac{15}{9} + \frac{5}{9} = 5.3$$

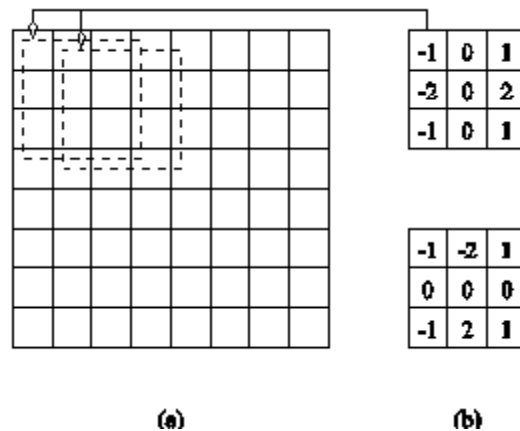
filtered image  $I'(x,y)$

0	0	0	0	0
0	5			0
0				0
0				0
0	0	0	0	0

# Parallel Time, Speedup, and Efficiency Example

Edge-detection problem requires us to apply a **3 x 3** template to each pixel.


If each multiply-add operation takes time  $t_c$ , the serial time for an  $n \times n$  image is given by  $T_s = 9 t_c n^2$ .



# Parallel Time, Speedup, and Efficiency

## Example (continued)

- One possible parallelization partitions the image equally into vertical segments, each with  $n^2 / p$  pixels.



- The boundary of each segment is  $2n$  pixels. This is also the number of pixel values that will have to be communicated. This takes time  $2(t_s + t_w n)$ .
- Templates may now be applied to all  $n^2 / p$  pixels in time  $T_s = 9 t_c n^2 / p$ .



# Parallel Time, Speedup, and Efficiency

## Example (continued)

- The total time for the algorithm is therefore given by:

$$T_P = 9t_c \frac{n^2}{p} + 2(t_s + t_w n)$$

- The corresponding values of speedup and efficiency are given by:

$$S = \frac{9t_c n^2}{9t_c \frac{n^2}{p} + 2(t_s + t_w n)}$$

$$E = \frac{1}{1 + \frac{2p(t_s + t_w n)}{9t_c n^2}}.$$

# Cost of a Parallel System

- Cost is the product of parallel runtime and the number of processing elements used ( $p T_p$ ).
- Cost reflects the sum of the time that each processing element spends solving the problem.
- A parallel system is said to be *cost-optimal* if the cost of solving a problem on a parallel computer is asymptotically identical to serial cost.
- Since  $E = T_s / p T_p$ , for cost optimal systems,  $E = O(1)$ .

# Cost of a Parallel System: Example

Consider the problem of adding  $n$  numbers on  $p$  processors.  
Assuming  $p = n$

- We have,  $T_p = \log n$
- The cost of this system is given by  $p T_p = n \log n$ .
- Since the serial runtime of this operation is  $\Theta(n)$ ,

$$E = \Theta(n/n \log n) = \Theta(1/\log n)$$

→ the algorithm is not cost optimal.

# Impact of Non-Cost Optimality

Consider a sorting algorithm that uses  $n$  processing elements to sort the list in time:

$$T_p = (\log n)^2$$

Serial runtime of a (comparison-based) sort is

$$T_s = n \log n$$

Then;

- Speedup:  $S = n / \log n$
- Efficiency:  $E = 1 / \log n$
- Cost  $C = n (\log n)^2$ .

This algorithm is not cost optimal but only by a factor of  $\log n$ .

If  $p < n$ , assigning  $n$  tasks to  $p$  processors gives:

- $T_p = n (\log n)^2 / p$ .
- $S = p / \log n$ .
- This speedup goes down as the problem size  $n$  is increased for a given  $p$  !

# Effect of Granularity on Performance

- Often, using fewer processors improves performance of parallel systems.
- Using fewer than the maximum possible number of processing elements to execute a parallel algorithm is called *scaling down* a parallel system.
- A naive way of scaling down is to think of each processor in the original case as a virtual processor and to assign virtual processors equally to scaled down processors.

# Amdahl's law

Limitations of inherent parallelism: a part  $s$  of the algorithm is not parallelizable

$$T_{seq} = (1-s).T_{seq} + s.T_{seq}$$



*parallelizable*   *not parallelizable*

$$T_{par} = \frac{(1-s).T_{seq}}{p} + s.T_{seq}$$

$$Speedup_{max} = \frac{T_{seq}}{T_{par}} = \frac{T_{seq}}{\frac{(1-s).T_{seq}}{p} + s.T_{seq}} = \frac{p}{1 + (p-1).s}$$

*Assume  
no other  
overhead*



$$\text{Speedup} < \frac{p}{1 + (p-1).s}$$

$$\text{Efficiency} < \frac{1}{1 + (p-1).s}$$

If  $p$  is big enough:


$$\text{Speedup} < \frac{1}{s}$$

<b>s</b>	<b>Speedup<sub>max</sub></b>
10%	10
25%	4
50%	2
75%	1.33

# Amdahl example: video decoding

Decoding 1080p video sequence

Stage	CPU (s)	CUDA (s)	
1 MOTION_DECODE	0.64	0.64	
2 MOTION_RENDER	16.16	1.33	← 12 ×
3 RESIDUAL_DECODE	12.00	12.94	
4 WAVELET_TRANSFORM	22.52	1.63	← 14 ×
5 COMBINE	11.27	0.39	← 29 ×
6 UPSAMPLE	14.53	0.85	← 17 ×
Total	77.13	17.76	← 4.3 ×






# Scalability of Parallel Systems

How do we extrapolate performance from small problems and small systems to larger problems on larger configurations?

Consider three parallel algorithms for computing an  $n$ -point Fast Fourier Transform (FFT) on 64 processing elements.



A comparison of the speedups obtained by the binary-exchange, 2-D transpose and 3-D transpose algorithms with  $t_c = 2$ ,  $t_w = 4$ ,  $t_s = 25$ , and  $t_h = 2$ .

Clearly, it is difficult to infer scaling characteristics from observations on small datasets on small machines.

# Scaling Characteristics of Parallel Programs

- The efficiency of a parallel program can be written as:

$$E = \frac{S}{p} = \frac{T_S}{pT_P}$$

or

$$E = \frac{1}{1 + \frac{T_o}{T_S}}.$$

- Derived from overhead function which is  $T_o = pT_p - T_s$
- The total overhead function  $T_o$  is an increasing function of  $p$ . This is because every program must contain some serial component. If this serial component of the program takes time  $t_{serial}$ , then during this time all the other processing elements must be idle. This corresponds to a total overhead function of  $(p-1)t_{serial}$ .

# Scaling Characteristics of Parallel Programs


- For a given problem size (i.e., the value of  $T_s$  remains constant), as we increase the number of processing elements,  $T_o$  increases.
- The overall efficiency of the parallel program goes down. This is the case for all parallel programs.

$$E = \frac{1}{1 + \frac{T_o}{T_s}}.$$

# Scaling Characteristics of Parallel Programs:

- Consider the problem of adding  $n$  numbers on  $p$  processing elements. Assume unit time for adding two numbers.
- We have seen that:

$$T_P = \frac{n}{p} + 2 \log p$$




$$S = \frac{n}{\frac{n}{p} + 2 \log p}$$

$$E = \frac{1}{1 + \frac{2p \log p}{n}}$$

The second phase involves  $\log p$  steps with a communication and an addition at each step. If a single communication takes unit time as well, the time for this phase is  $2 \log p$ .

# Scaling Characteristics of Parallel Programs: Example (continued)

Plotting the speedup for various input sizes gives us:



- Speedup versus the number of processing elements for adding a list of numbers.
- Speedup tends to saturate and efficiency drops as a consequence of Amdahl's law.
- A larger instance of the same problem yields higher speedup and efficiency for the same number of processing elements, although both speedup and efficiency continue to drop with increasing  $p$ .

# Scaling Characteristics of Parallel Programs

- Total overhead function  $T_o$  is a function of both problem size  $T_s$  and the number of processing elements  $p$ . In many cases,  $T_o$  grows sublinearly with respect to  $T_s$ .
- In such cases, the efficiency increases if the problem size is increased keeping the number of processing elements constant.
- For such systems, we can simultaneously increase the problem size and number of processors to keep efficiency constant.
- We call such systems *scalable parallel systems*.

n	p=1	p=4	p=8	p=16	p=32
64	1	0.8	0.57	0.33	0.17
192	1	0.92	0.8	0.6	0.38
320	1	0.95	0.87	0.71	0.5
512	1	0.97	0.91	0.8	0.62

# Scaling Characteristics of Parallel Programs



- Recall that cost-optimal parallel systems have an efficiency of  $\Theta(1)$ .
- Scalability and cost-optimality are therefore related.
- A scalable parallel system can always be made cost-optimal if the number of processing elements and the size of the computation are chosen appropriately.

# Isoefficiency Metric of Scalability

- For a given problem size, as we increase the number of processing elements, the overall efficiency of the parallel system goes down for all systems.
- For some systems, the efficiency of a parallel system increases if the problem size is increased while keeping the number of processing elements constant.



# Isoefficiency Metric of Scalability



Variation of efficiency:

- (a) as the number of processing elements is increased for a given problem size;
- (b) as the problem size is increased for a given number of processing elements.

-The phenomenon illustrated in graph (b) is not common to all parallel systems-

# Isoefficiency Metric of Scalability

- What is the rate at which the problem size must increase with respect to the number of processing elements to keep the efficiency fixed?
- This rate determines the scalability of the system.
- Before we formalize this rate, we define the problem size  **$W$**  as the number of basic computation steps in the best serial algorithm to solve the problem on a single processing element.

# Isoefficiency Metric of Scalability

- We can write parallel runtime as:

$$T_P = \frac{W + T_o(W, p)}{p}$$

The resulting expression for speedup is

$$\begin{aligned} S &= \frac{W}{T_P} \\ &= \frac{Wp}{W + T_o(W, p)}. \end{aligned}$$

- Finally, we write the expression for efficiency as

$$\begin{aligned} E &= \frac{S}{p} \\ &= \frac{W}{W + T_o(W, p)} \\ &= \frac{1}{1 + T_o(W, p)/W}. \end{aligned}$$

# Isoefficiency Metric of Scalability

- For scalable parallel systems, efficiency can be maintained at a fixed value if the ratio  $T_o / W$  is maintained at a constant value.
- For a desired value  $E$  of efficiency,

$$E = \frac{1}{1 + T_o(W, p)/W},$$
$$\frac{T_o(W, p)}{W} = \frac{1 - E}{E},$$
$$W = \frac{E}{1 - E} T_o(W, p).$$

- If  $K = E / (1 - E)$  is a constant depending on the efficiency to be maintained, since  $T_o$  is a function of  $W$  and  $p$ , we have

$$W = K T_o(W, p).$$

# Isoefficiency Metric of Scalability

- The problem size  $W$  can usually be obtained as a function of  $p$  by algebraic manipulations to keep efficiency constant.
- This function is called the *isoefficiency function*.
- This function determines the ease with which a parallel system can maintain a constant efficiency and hence achieve speedups increasing in proportion to the number of processing elements.

# Isoefficiency Metric: Example

- The overhead function for the problem of adding  $n$  numbers on  $p$  processing elements is approximately  $2p \log p$ .
- Substituting  $T_o$  by  $2p \log p$ , we get

$$W = K 2p \log p.$$

Thus, the asymptotic isoefficiency function for this parallel system is

$$\Theta(p \log p)$$

- If the number of processing elements is increased from  $p$  to  $p'$ , the problem size must be increased by a factor of  $(p' \log p') / (p \log p)$  to get the same efficiency as on  $p$  processing elements.

# Reading List

- “Introduction to Parallel Computing”, 2nd Edition, 2003, Addison Wesley
  - By Ananth Grama; Anshul Gupta; George Karypis; Vipin Kumar
- Chapter 5: Analytical Modeling of Parallel Systems
  - <http://proquestcombo.safaribooksonline.com/0201648652/ch05>

# Performance Analysis

Slides adapted from:

Parallel Systems: Performance Analysis of Parallel Processing

PhD Thesis, Jan Lemeire

November 6, 2007






# Goals of Performance Analysis


- Understanding of the computational process in terms of resource consumption
- Identification of inefficient patterns
- Performance prediction
- Performance characterization of program and system

# Speedup vs. # of processors




- 1) Ideal, linear speedup
- 2) Increasing, sub-linear speedup
- 3) Speedup with an optimal number of processors
- 4) No speedup
- 5) Super-linear speedup

# Speedup vs. problem size



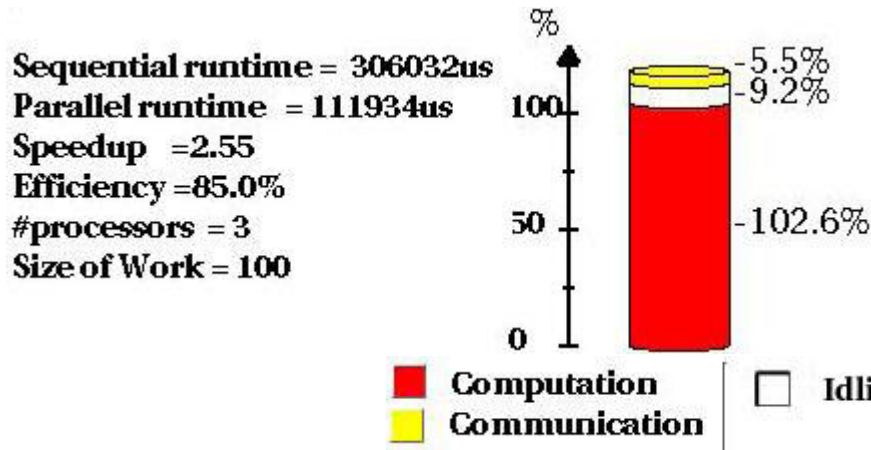
- 1) Constant speedup
- 2) Increasing, asymptotically, towards value sublinear speedup ( $< p$ )
- 3) Increasing towards  $p$
- 4) Increasing towards super-linear speedup

# Parallel Matrix Multiplication



Speedup=2.55 Efficiency = 85%


# Parallel Matrix Multiplication



Parallel anomaly = 2.6%

- Overheads: the communication and the idle time.
- Their ratio with the sequential time is given.
- The sum of the processor's computation times divided by the sequential runtime is also given, but is not equal to 100%. A value of 100% means that the computation time of the useful work is equal for a sequential as for a parallel execution.
- It is 102.6% instead, which means that the overhead ratio of the parallel anomaly is 2.6%. In parallel, 2.6% more cycles are needed to do the same work.

# Parallel Matrix Multiplication


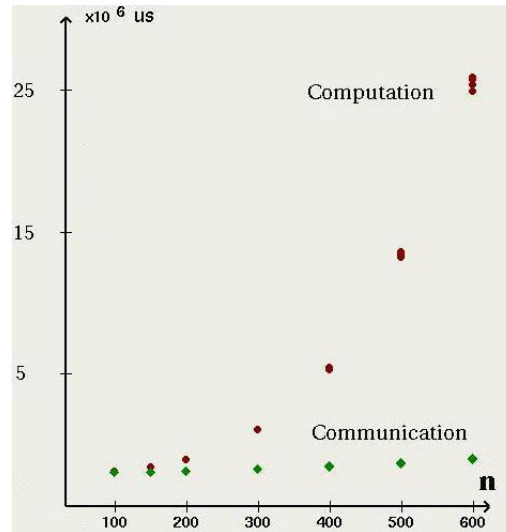
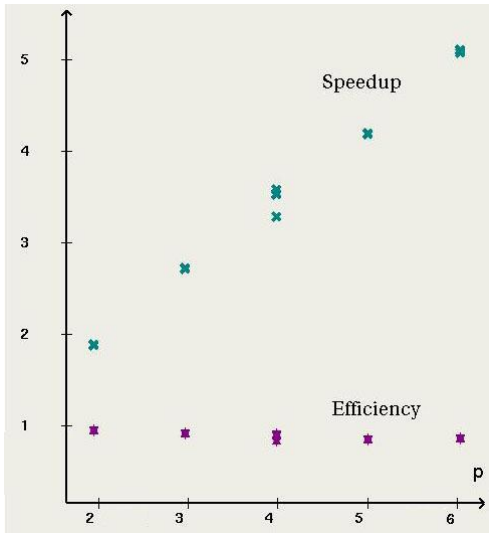


# Overhead Classification

- Control of parallelism: extra functionality necessary for parallelization (like partitioning)
- Communication: overhead time not overlapping with computation
- Idling: processor has to wait for further information
- Parallel anomaly : useful work differs for sequential and parallel execution

$$T_{seq} + T_{anomaly} = \sum_i^p T_{work}^i$$

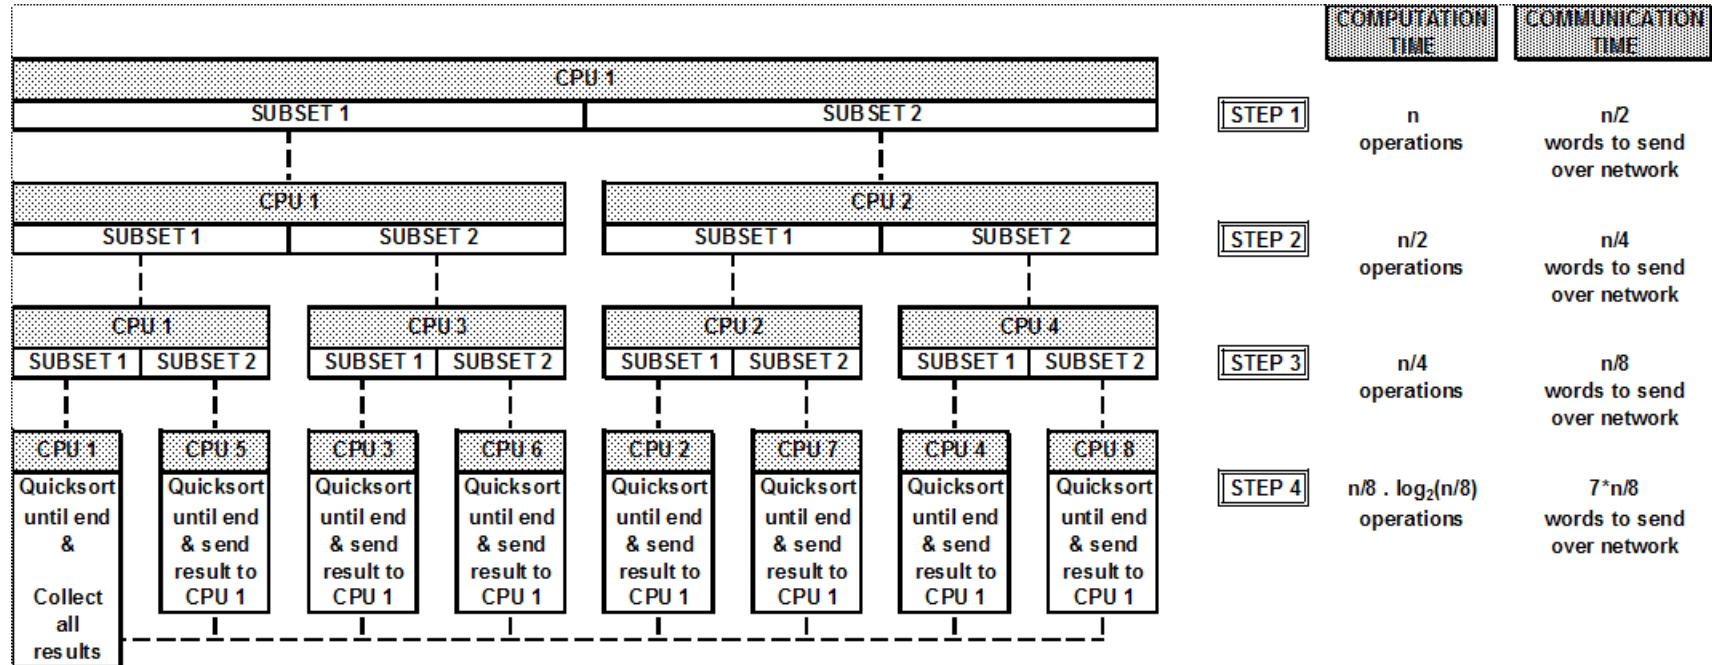
# Overhead Classification



$P$ : no. of processors  
 $n$ : work size (matrix size)



# Quicksort



# Overhead Optimization

1. Generate/draw execution profile
2. Identify lost cycles
3. Study impact on overhead
4. Determine causes of overhead
5. Plot performance in function of  $p$  and  $W$