

# Thrust Library

Dr. Alptekin Temizel  
atemizel@metu.edu.tr



# Thrust Library

- Distributed with CUDA Toolkit
- Header-only library
- Architecture agnostic
- Just compile and run!

# Thrust Objectives

- Programmer productivity
  - Rapidly develop complex applications
  - Leverage parallel primitives
- Encourage generic programming
  - Don't reinvent the wheel
- High performance
  - With minimal programmer effort
- Interoperability
  - Integrates with CUDA C/C++ code

# Outline

- Thrust Basics

- Containers: `host_vector`, `device_vector`
- Host – device copy operations
- Initialization and modification: `fill`, `sequence`, `generate`, `random`
- Device and raw pointers
- STL integration
- Templates and functors
- Modifying vectors: `replace`, `replace_if`, `reverse`, `reverse_copy`
- Namespaces

- Algorithms

- Element wise: `for_each`, `transform`
- Reduction: `reduce`, `reduce_by_key`, `inner_product`, `find`, `equal` ...
- Kernel fusion and `transform_reduce`
- Iterators: `constant_iterator`, `counting_iterator`, `zip_iterator`
- Prefix-sums: `inclusive_scan`, `inclusive_scan_by_key`
- Sorting: `sort`, `stable_sort`, `sort_by_key`

# Containers

- C++ template library for CUDA
  - Mimics Standard Template Library (STL)
- Make common operations concise and readable
  - Hides cudaMalloc, cudaMemcpy and cudaFree
- Containers
  - thrust::host\_vector<T>
    - CPU implementation via OpenMP
  - thrust::device\_vector<T>
    - GPU implementation via CUDA
  - <T> stands for a generic data type

# Thrust Algorithms

Algorithm	Description
<code>reduce</code>	Sum of a sequence
<code>find</code>	First position of a value in a sequence
<code>mismatch</code>	First position where two sequences differ
<code>inner_product</code>	Dot product of two sequences
<code>equal</code>	Whether two sequences are equal
<code>min_element</code>	Position of the smallest value
<code>count</code>	Number of instances of a value
<code>is_sorted</code>	Whether sequence is in sorted order
<code>transform_reduce</code>	Sum of transformed sequence

# Example

```
#include <thrust/host_vector.h>
#include <thrust/device_vector.h>
#include <thrust/sort.h>
#include <cstdlib>

int main(void)
{
    // generate 32M random numbers on the host
    thrust::host_vector<int> h_vec(32 << 20);
    thrust::generate(h_vec.begin(), h_vec.end(), rand);


    // transfer data to the device
    thrust::device_vector<int> d_vec = h_vec;

    // sort data on the device
    thrust::sort(d_vec.begin(), d_vec.end());

    // transfer data back to host
    thrust::copy(d_vec.begin(), d_vec.end(), h_vec.begin());

    return 0;
}
```

# Interoperability





# Backend System Options

## Host Systems

**THRUST\_HOST\_SYSTEM\_CPP**

**THRUST\_HOST\_SYSTEM\_OMP**

**THRUST\_HOST\_SYSTEM\_TBB**

OpenMP

Intel Threading  
Building Blocks

## Device Systems

**THRUST\_DEVICE\_SYSTEM\_CUDA**

**THRUST\_DEVICE\_SYSTEM\_OMP**

**THRUST\_DEVICE\_SYSTEM\_TBB**

# Multiple Backend Systems

- Mix different backends freely within the same app

```
thrust::omp::vector<float> my_omp_vec(100);  
thrust::cuda::vector<float> my_cuda_vec(100);  
  
...  
  
// reduce in parallel on the CPU  
thrust::reduce(my_omp_vec.begin(), my_omp_vec.end());  
  
// sort in parallel on the GPU  
thrust::sort(my_cuda_vec.begin(), my_cuda_vec.end());
```

# Host ↔ Device Copy Operations

```
#include <thrust/host_vector.h>
#include <thrust/device_vector.h>
#include <iostream>

// allocate host vector with 3 elements
thrust::host_vector<int> h_vec(3);
h_vec[0] = 13;           // array operator is overloaded
h_vec[1] = 27;
h_vec[2] = -36;
std::cout << "h_vec has size " << h_vec.size() << std::endl;


// copy host vector to device (via copy constructor)
thrust::device_vector<int> d_vec = h_vec;

// manipulate device values from the host
d_vec[2] = 36;           // not efficient
```

# Host ↔ Device Copy Operations

```
#include <thrust/copy.h>
```

```
// copy all of d_vec back to the beginning of h_vec  
thrust::copy(d_vec.begin(), d_vec.end(), h_vec.begin());
```



```
// [begin, end) pair defines a sequence of N elements  
// N equals d_vec.size(), same as:  
thrust::device_vector<int>::iterator begin = d_vec.begin();  
thrust::device_vector<int>::iterator end   = d_vec.end();  
// compute size of sequence [begin, end)  
int N = end - begin;
```

# Basics of Iterators

- Iterators act like pointers

```
// declare iterator variables
device_vector<int>::iterator begin = d_vec.begin();
device_vector<int>::iterator end   = d_vec.end();

// pointer arithmetic
begin++;

// dereference device iterators from the host
int a = *begin;
int b = begin[3];

// compute size of range [begin,end)
int size = end - begin;
```

# Vector Initialization

```
#include <thrust/fill.h>
#include <thrust/sequence.h>
#include <thrust/generate.h>

// initialize all ten integers of a device_vector to 1
thrust::device_vector<int> d_vec(10, 1);

// set the first seven elements of a vector to 9
thrust::fill(d_vec.begin(), d_vec.begin() + 7, 9);

// set the elements of h_vec to 0, 1, 2, 3, ...
thrust::sequence(h_vec.begin(), h_vec.end());

// generate 16M random numbers on the host
thrust::host_vector<int> h_vec(1 << 24);
thrust::generate(h_vec.begin(), h_vec.end(), rand);
```

# Device and Raw Pointers

```
size_t N = 10;
```

```
int* dev_raw_ptr1;
```

```
cudaMalloc((void **)&dev_raw_ptr1, N * sizeof(int));
```

```
thrust::device_ptr<int> dev_ptr(dev_raw_ptr1 );
```

```
// use device_ptr in thrust algorithms
```

```
thrust::fill(dev_ptr, dev_ptr + N, (int)0);
```

# Device and Raw Pointers

```
size_t N = 10;
```

```
thrust::device_vector<int> d_vec(N);
```

```
int* dev_raw_ptr2 = thrust::raw_pointer_cast(&d_vec[0]);
```



# Device and Raw Pointers

```
size_t N = 10;
```

```
thrust::device_ptr<int> dev_ptr = thrust::device_malloc<int>(N);  
int* dev_raw_ptr3 = thrust::raw_pointer_cast( dev_ptr );
```

# Interoperability

- Convert iterators to raw pointers

```
// allocate device vector
const int N = 65536;
thrust::device_vector<int> d_vec(N);

// obtain raw pointer to device vector's memory
int* dev_raw_ptr = thrust::raw_pointer_cast(&d_vec[0]);

// use dev_raw_ptr in a CUDA C kernel

const int number = 29;
addNumber<<< N / 256, 256 >>>(N, dev_raw_ptr, number);
```

---

```
__global__ void addNumber(int N, int* data, int number)
{
    int i = blockIdx.x * 256 + threadIdx.x;
    if (i < N) {
        data[i] += number;
    }
}
```

# STL Integration

- Compatible with STL containers
  - Eases integration for vector, list, map, ...

```
// list container on host
```

```
std::list<int> h_list;
```

```
h_list.push_back(13);
```

```
h_list.push_back(27);
```

```
h_list.push_back(36);
```

```
// approach 1: create a device vector and copy an STL list into  
a device_vector
```

```
thrust::device_vector<int> d_vec(h_list.size());
```

```
thrust::copy(h_list.begin(), h_list.end(), d_vec.begin());
```

```
// approach 2: initialize a device_vector with the list
```

```
thrust::device_vector<int> d_vec(h_list.begin(), h_list.end());
```



# Templates and Functors - I

- Function templates

```
// function template to add numbers (type of T is variable)
```

```
template<typename T>
```

```
T add(T a, T b)
```

```
{
```

```
    return a + b;
```

```
}
```

```
// add integers
```

```
int x = 10; int y = 20; int z;
```

```
z = add<int>(x,y);           // type of T explicitly specified
```

```
z = add(x,y);               // type of T determined automatically
```

```
// add floats
```

```
float x = 10.0f; float y = 20.0f; float z;
```

```
z = add<float>(x,y);         // type of T explicitly specified
```

```
z = add(x,y);               // type of T determined automatically
```

# Templates and Functors - II

- Function objects (Functors)

```
// templated functor to add numbers
template<typename T>
class add
{
    public:
    T operator() (T a, T b)
    {
        return a + b;
    }
};

int x = 10; int y = 20; int z;
add<int> func;    // create an add functor for T=int
z = func(x,y);   // invoke functor on x and y

float x = 10; float y = 20; float z;
add<float> func;  // create an add functor for T=float
z = func(x,y);   // invoke functor on x and y
```

# Templates and Functors - III

- Generic Algorithms

```
// apply function f to sequences x, y and store result in z
template <typename T, typename Function>
void transform(int N, T* x, T* y, T* z, Function f)
{
    for (int i = 0; i < N; i++) {
        z[i] = f(x[i], y[i]);
    }
}

int N = 100;
int x[N]; int y[N]; int z[N];

add<int> func; // add functor for T=int
transform(N, x, y, z, func); // compute z[i] = x[i] + y[i]

transform(N, x, y, z, add<int>() ); // equivalent
```

# Modifying Vectors - I

```
#include <thrust/replace.h>

// create a device vector using a host array
const int N = 7;
int arr[N] = {1, 7, 1, 1, 3, 2, -1};
thrust::device_vector<int> d_vec(arr, arr + N) ;

// replace all the ones in d_vec with tens
thrust::replace(d_vec.begin(), d_vec.end(), 1, 10);

// replace via a predicate (E.g. greater than a number)
struct is_greater_than {
    int threshold;

    is_greater_than(int t) { threshold = t; }

    __host__ __device__
    bool operator()(int x) { return x > threshold; }
};

is_greater_than pred(2);
thrust::replace_if(d_vec.begin(), d_vec.end(), pred, 10);
```

# Modifying Vectors - II

- Reversing a vector

```
#include <thrust/reverse.h>
```

```
const int N = 6;
```

```
int data[N] = {0, 1, 2, 3, 4, 5};
```

```
thrust::device_vector<int> dev_vec(data, data + N);
```

```
// takes reverse of a vector inplace
```

```
thrust::reverse(dev_vec.begin(), dev_vec.end());
```

```
// dev_vec is now {5, 4, 3, 2, 1, 0}
```

```
// reverse_copy: reversed is written to a different output
```

```
thrust::device_vector<int> output(N);
```

```
thrust::reverse_copy(dev_vec.begin(), dev_vec.end(),  
                    output.begin());
```



# Namespaces

- Avoid name collisions

```
// allocate host memory
thrust::host_vector<int> h_vec(10);

// call STL sort
std::sort(h_vec.begin(), h_vec.end());

// call Thrust sort
thrust::sort(h_vec.begin(), h_vec.end());

// for brevity
using namespace thrust;

// without namespace
int sum = reduce(h_vec.begin(), h_vec.end());
```

# Algorithms

- Element wise
  - for\_each, transform
- Reduction
  - reduce, reduce\_by\_key, inner\_product, find, equal ...
- Kernel fusion
  - transform\_reduce
- Iterators
  - constant\_iterator, counting\_iterator, zip\_iterator
- Prefix-sums
  - inclusive\_scan, inclusive\_scan\_by\_key
- Sorting
  - sort, stable\_sort, sort\_by\_key

# Algorithms

- Thrust targets maximal occupancy and will compare the resource usage of the kernel (e.g., number of registers, amount of shared memory) with the resources of the target GPU to determine a launch configuration with highest occupancy.
- While the maximal occupancy heuristic is not necessarily optimal, it is straightforward to compute and effective in practice.

# Element wise : for\_each

```
struct normalize_functor
: public unary_function<double4, double4> {
__device__ __host__ double4 operator()(double4 v)
{
    double len = sqrt(v.x*v.x + v.y*v.y + v.z*v.z);
    v.x /= len;
    v.y /= len;
    v.z /= len;
}
}

device_vector<double4> d_vec = ...;
for_each(d_vec.begin(), d_vec.end(), normalize_functor() );
```

# Element wise : transform

```
#include <thrust/transform.h>

// allocate memory
device_vector<int> A(10);
device_vector<int> B(10);
device_vector<int> C(10);

// transform A + B -> C
transform(A.begin(), A.end(),
          B.begin(), C.begin(),
          plus<int>());

// transform A - B -> C
transform(A.begin(), A.end(),
          B.begin(), C.begin(),
          minus<int>());
```

# SAXPY via CUDA kernel

---

```
__global__  
void saxpy_kernel(int n, float a, float * x, float * y)  
{  
    const int i = blockDim.x * blockIdx.x + threadIdx.x;  
  
    if (i < n)  
        y[i] = a * x[i] + y[i];  
}  
  
void saxpy(int n, float a, float * x, float * y)  
{  
    // set launch configuration parameters  
    int block_size = 256;  
    int grid_size = (n + block_size - 1) / block_size;  
  
    // launch saxpy kernel  
    saxpy_kernel<<< grid_size, block_size >>>(n, a, x, y);  
}
```

---

(a) CUDA C

# SAXPY via Thrust transform

---

```
struct saxpy_functor
{
    const float a;

    saxpy_functor(float _a) : a(_a) {}

    __host__ __device__
    float operator()(float x, float y)
    {
        return a * x + y;
    }
};

void saxpy(float a, device_vector<float>& x, device_vector<float>& y)
{
    // setup functor
    saxpy_functor func(a);

    // call transform
    transform(x.begin(), x.end(), y.begin(), y.begin(), func);
}
```

---

(b) Thrust

# Element wise : custom transform

```
struct negate_float2 {  
    __host__ __device__  
    float2 operator()(float2a)  
    {  
        return make_float2(-a.x, -a.y);  
    }  
};  
  
// declare storage  
device_vector<float2> input = ...  
device_vector<float2> output = ...  
  
// create function object or 'functor'  
negate_float2 func;  
  
// negate vectors  
transform(input.begin(), input.end(),  
          output.begin(), func);
```



# Reduction

Algorithm	Description
<code>reduce</code>	Sum of a sequence
<code>find</code>	First position of a value in a sequence
<code>mismatch</code>	First position where two sequences differ
<code>inner_product</code>	Dot product of two sequences
<code>equal</code>	Whether two sequences are equal
<code>min_element</code>	Position of the smallest value
<code>count</code>	Number of instances of a value
<code>is_sorted</code>	Whether sequence is in sorted order
<code>transform_reduce</code>	Sum of transformed sequence

# Reduction : reduce

```
#include <thrust/reduce.h>

// declare storage
host_vector<int>      i_vec = ...
device_vector<float> f_vec = ...

// sum of integers (equivalent calls)
reduce(i_vec.begin(), i_vec.end());
reduce(i_vec.begin(), i_vec.end(), 0, plus<int>());

// sum of floats (equivalent calls)
reduce(f_vec.begin(), f_vec.end());
reduce(f_vec.begin(), f_vec.end(), 0.0f, plus<float>());

// maximum of integers
reduce(i_vec.begin(), i_vec.end(), 0, maximum<int>());
```

# Reduction : reduce – extrema calculation

```
template<typename ForwardIterator >
```

```
ForwardIterator thrust::min_element (ForwardIterator first, ForwardIterator last)
```

```
template<typename ForwardIterator , typename BinaryPredicate >
```

```
ForwardIterator  
thrust::min_element (ForwardIterator first, ForwardIterator last, BinaryPredicate comp)
```

```
template<typename ForwardIterator >
```

```
ForwardIterator thrust::max_element (ForwardIterator first, ForwardIterator last)
```

```
template<typename ForwardIterator , typename BinaryPredicate >
```

```
ForwardIterator  
thrust::max_element (ForwardIterator first, ForwardIterator last, BinaryPredicate comp)
```

```
template<typename ForwardIterator >
```

```
thrust::pair< ForwardIterator, thrust::minmax_element (ForwardIterator first, ForwardIterator last)
```

```
template<typename ForwardIterator , typename BinaryPredicate >
```

```
thrust::pair< ForwardIterator, ForwardIterator > thrust::minmax_element (ForwardIterator first, ForwardIterator last, BinaryPredicate comp)
```

# Reduction : count\_if

```
#include <thrust/count_if.h>

// count via a predicate (E.g. greater than a number)
struct is_greater_than {
    int threshold;

    is_greater_than(int t) { threshold = t; }

    __host__ __device__
    bool operator()(int x) { return x > threshold; }
};

// create a device vector using a host array
const int N = 10;
int arr[N] = {1, 7, 1, 1, 3, 2, -1, 12, 6, 32};
thrust::device_vector<int> d_vec(arr, arr + N) ;

is_greater_than pred(2);
int count = thrust::count_if(d_vec.begin(), d_vec.end(), pred);
```

# Iterators

- Behave like pointers
- Keep track of memory spaces
- Convertible to raw pointers
- Examples
  - `constant_iterator`,
  - `counting_iterator`,
  - `transform_iterator`
  - `permutation_iterator`
  - `zip_iterator`

# Iterators - constant\_iterator

- Mimics an infinite array filled with a constant value

```
// create iterators
constant_iterator<int> begin(10);
constant_iterator<int> end = begin + 3;

begin[0]    // returns 10
begin[1]    // returns 10
begin[100]  // returns 10

// sum of [begin, end)
reduce(begin, end);    // returns 30 (i.e. 3 * 10)
```




# Iterators - counting\_iterator

- Mimics an infinite array with sequential values

```
// create iterators
counting_iterator<int> begin(10);
counting_iterator<int> end = begin + 3;

begin[0]    // returns 10
begin[1]    // returns 11
begin[100]  // returns 110

// sum of [begin, end)
reduce(begin, end);    // returns 33 (i.e. 10 + 11 + 12)
```



# Iterators - transform\_iterator


- Conserves memory capacity and bandwidth

```
// initialize vector
device_vector<int> vec(3);
vec[0] = 10; vec[1] = 20; vec[2] = 30;

// create iterator (type omitted)
begin = make_transform_iterator(vec.begin(), negate<int>());
end    = make_transform_iterator(vec.end(),    negate<int>());

begin[0]    // returns -10
begin[1]    // returns -20
begin[2]    // returns -30

// sum of [begin, end)
reduce(begin, end);    // returns -60 (i.e. -10 + -20 + -30)
```





# Iterators – zip\_iterator


- Support coalesced access via Structure of Arrays (SoA)

```
// initialize vectors
device_vector<int>  A(3);
device_vector<char> B(3);
A[0] = 10;  A[1] = 20;  A[2] = 30;
B[0] = 'x'; B[1] = 'y'; B[2] = 'z';

// create iterator (type omitted)
begin = make_zip_iterator(make_tuple(A.begin(), B.begin()));
end    = make_zip_iterator(make_tuple(A.end(),   B.end()));

begin[0] // returns tuple(10, 'x')
begin[1] // returns tuple(20, 'y')
begin[2] // returns tuple(30, 'z')

// maximum of [begin, end)
maximum< tuple<int,char> > binary_op;
reduce(begin, end, begin[0], binary_op); // returns tuple(30, 'z')
```



# Iterators - Slow Vector Rotation

```
struct rotate_float3{
    __host__ __device__
    float3 operator()(float3 v) {
        float x = v.x;
        float y = v.y;
        float z = v.z;
        float rx= 0.36f*x + 0.48f*y - 0.80f*z;
        float ry=-0.80f*x + 0.60f*y + 0.00f*z;
        float rz= 0.48f*x + 0.64f*y + 0.60f*z;

        return make_float3(rx, ry, rz);
    }
};

device_vector<float3> v(n);
transform(v.begin(), v.end(), v.begin(), rotate_float3() );
```



# Iterators - Fast Vector Rotation

```
struct rotate_tuple{
    __host__ __device__
    tuple<float,float,float> operator() (tuple<float,float,float> v) {
        float x = get<0>(v);
        float y = get<1>(v);
        float z = get<2>(v);
        float rx= 0.36f*x + 0.48f*y - 0.80f*z;
        float ry=-0.80f*x + 0.60f*y + 0.00f*z;
        float rz= 0.48f*x + 0.64f*y + 0.60f*z;
        return make_tuple(rx, ry, rz);
    }
};
```

```
device_vector<float> x(n), y(n), z(n);
transform(
    make_zip_iterator(make_tuple(x.begin(), y.begin(), z.begin())),
    make_zip_iterator(make_tuple(x.end(), y.end(), z.end())),
    rotate_tuple() );
```



# Kernel fusion - I

- Consider  $z = g(f(x))$

```
device_vector<float> x(n);           // input
device_vector<float> f_x(n);         // temporary
device_vector<float> z(n);           // result

// compute f(x)
transform(x.begin(), x.end(), f_x.begin(), f());

// compute g(f(x))
transform(f_x.begin(), f_x.end(), z.begin(), g());
```

- Storage:  $3 * n$
- Bandwidth:  $2 * n$  reads +  $2 * n$  writes
- Temporaries:  $n$

# Kernel fusion - II

- A better way with **transform\_iterator**

```
device_vector<float> x(n);      // input
device_vector<float> z(n);      // result

// compute g(f(x))
transform(make_transform_iterator(x.begin(), f()),
         make_transform_iterator(x.end(), f()),
         z.begin(), g());
```

- **Storage:  $2 * n$**
- **Bandwidth:  $n$  reads +  $n$  writes**
- **Temporaries: 0**

# Fusion: Sum of squares $\sum x_i^2$

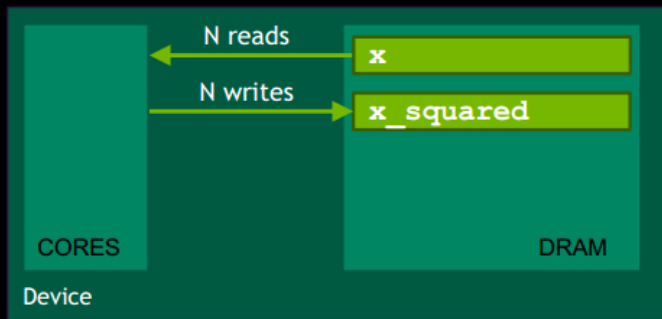
```
struct square { __device__ __host__ float operator()(float xi) { return xi*xi; } };

float sum_of_squares(const thrust::device_vector<float> &x)
{
    size_t N = x.size();
    thrust::device_vector<float> x_squared(N); // Temporary storage: N elements.

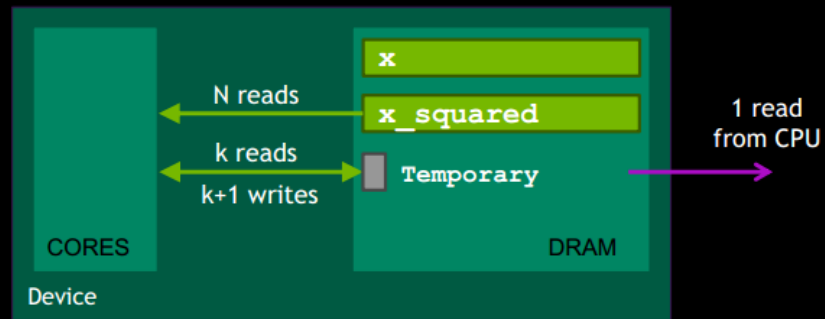
    // Compute x^2: N reads + N writes.
    thrust::transform(x.begin(), x.end(), x_squared.begin(), square());

    // Compute the sum of x^2s: N + k reads + k+1 writes (k is a small constant).
    return thrust::reduce(x_squared.begin(), x_squared.end());
}
```

transform



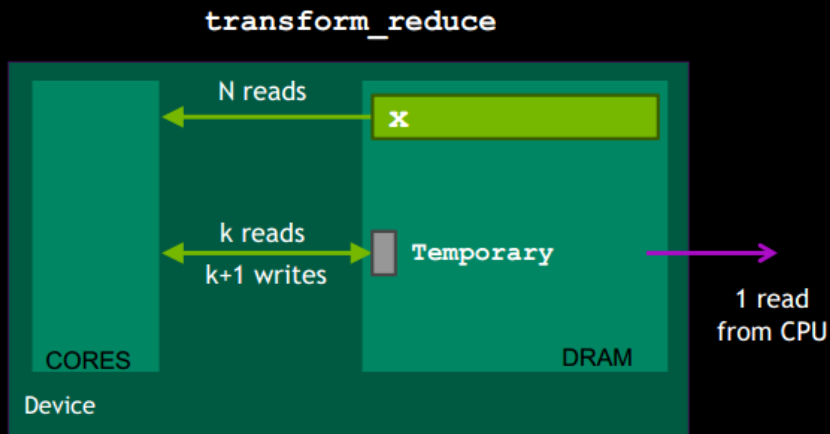
reduce



# Fusion

- Combined related operations together

```
float fused_sum_of_squares(const thrust::device_vector<float> &x)
{
    // Compute the x^2s and their sum: N + k reads + k+1 writes (k is a small constant).
    return thrust::reduce(
        thrust::make_transform_iterator(x.begin(), square()),
        thrust::make_transform_iterator(x.end(), square()));
}
```



We save:

- N temporary storage ( $x\_squared$ )
- N writes (to  $x\_squared$ )
- N reads (from  $x\_squared$ )

# Kernel fusion - III

- Slow Vector Norm

```
// define transformation f(x) -> x^2
struct square{
    __host__ __device__
    float operator()(float x)
    {
        return x * x;
    }
};

device_vector<float> x_2(n); // temporary storage
transform(x.begin(), x.end(), x_2.begin(), square());
return sqrt(reduce(x_2.begin(), x_2.end(),
                  0.0f, plus<float>()));
```



# Kernel fusion - IV

- Fast Vector Norm by Fusion

```
// define transformation f(x) -> x^2
```

```
struct square{  
    __host__ __device__  
    float operator() (float x)  
    {  
        return x * x;  
    }  
};
```

```
// fusion with transform_iterator
```

```
return sqrt(  
    reduce(make_transform_iterator(x.begin(), square()),  
        make_transform_iterator(x.end(), square()),  
        0.0f, plus<float>()) );
```

# Kernel fusion - V

- Fast Vector Norm using transform\_reduce

```
// define transformation f(x) -> x^2
```

```
struct square{  
    __host__ __device__  
    float operator()(float x)  
    {  
        return x * x;  
    }  
};
```

**Speedup:**  
**7.0x (GTX 480)**  
**4.4x (GTX 280)**

```
// fusion with transform_reduce
```

```
return sqrt(  
    transform_reduce(x.begin(), x.end(),  
                    square()),  
    0.0f, plus<float>()));
```

# Prefix-Sums / Scan Operations

- Important building blocks in many parallel algorithms
  - such as stream compaction and radix sort
  - `transform_inclusive_scan` and `transform_exclusive_scan` applies a unary function before scan operation

```
// inclusive scan operation using the default plus operator
#include <thrust/scan.h>
```

```
int data [6] = {1, 0, 2, 2, 1, 3};
```

```
inclusive_scan(data , data + 6, data ); // in - place scan
// data is now {1, 1, 3, 5, 6, 9}
```

```
exclusive_scan(data , data + 6, data ); // in - place scan
// data is now {0, 1, 1, 3, 5, 6}
```





# Sorting : sort

```
// generate 16M random numbers on the host
thrust::host_vector<int> h_vec(1 << 24);
thrust::generate(h_vec.begin(), h_vec.end(), rand);

// transfer data to the device
thrust::device_vector<int> d_vec = h_vec;

// sort data on the device
thrust::sort(d_vec.begin(), d_vec.end());
```



# Sorting : stable\_sort, sort\_by\_key

- Stable sort preserves the relative ordering of equivalent elements

```
thrust::stable_sort(d_vec.begin(), d_vec.end());
```

- Sorts the data by given key

```
const int N = 6;  
int keys[N] = { 1, 4, 2, 8, 5, 7};  
char values[N] = {'a', 'b', 'c', 'd', 'e', 'f'};  
thrust::sort_by_key(keys, keys + N, values );  
// keys is now { 1, 2, 4, 5, 7, 8}  
// values is now {'a ', 'c ', 'b ', 'e ', 'f ', 'd '}
```

# Sources

- Nvidia Research

- “CUDA Toolkit 4.0 - Thrust Quick Start Guide”
- “An Introduction To Thrust”, Jared Hoberock & Nathan Bell
- “High-Productivity CUDA Development with the Thrust Template Library”, Nathan Bell
- “Thrust by Example - Advanced Features and Techniques”, Jared Hoberock
- “Rapid Problem Solving using Thrust”, Nathan Bell
- “Thrust: A Productivity-Oriented Library for CUDA”, Nathan Bell and Jared Hoberock

- Others

- “CUDA Tricks”, Damodaran Ramani