

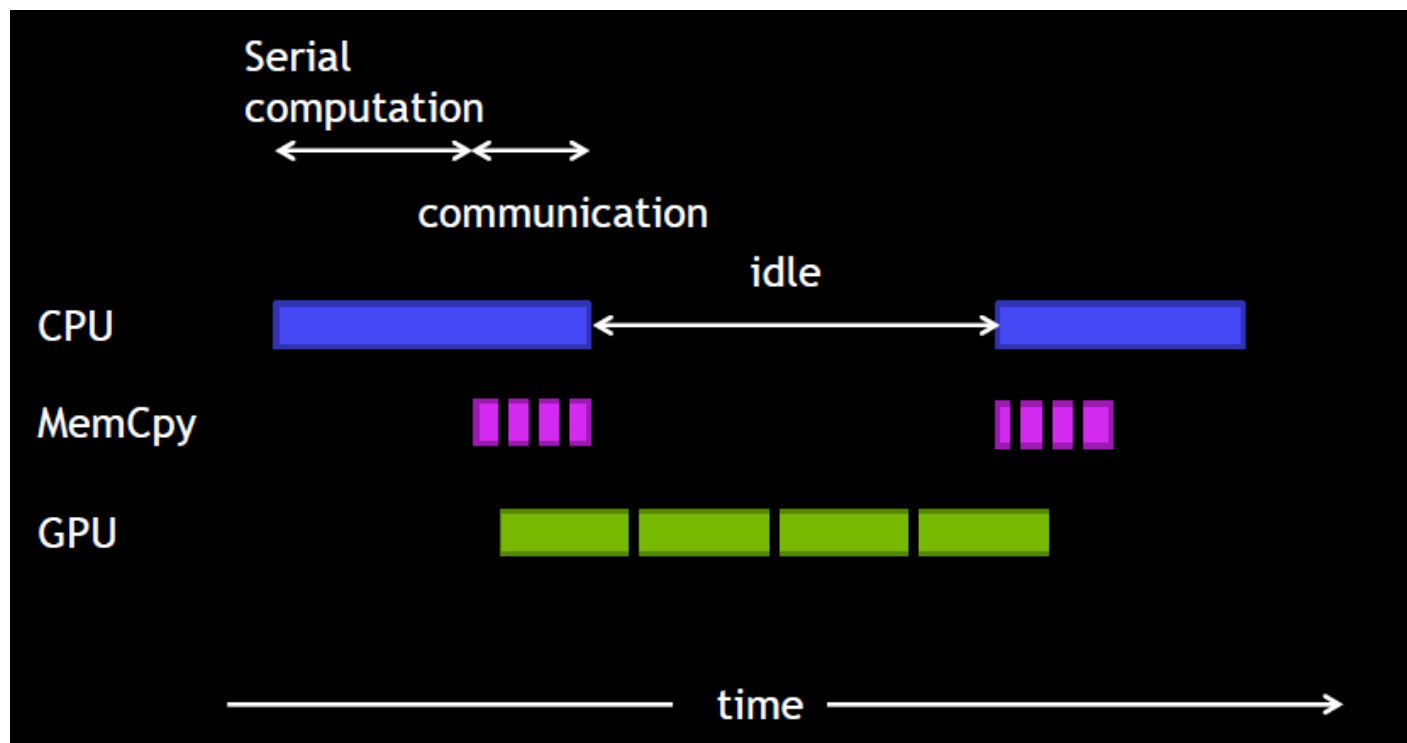



ORTA DOĞU TEKNİK ÜNİVERSİTESİ
MIDDLE EAST TECHNICAL UNIVERSITY

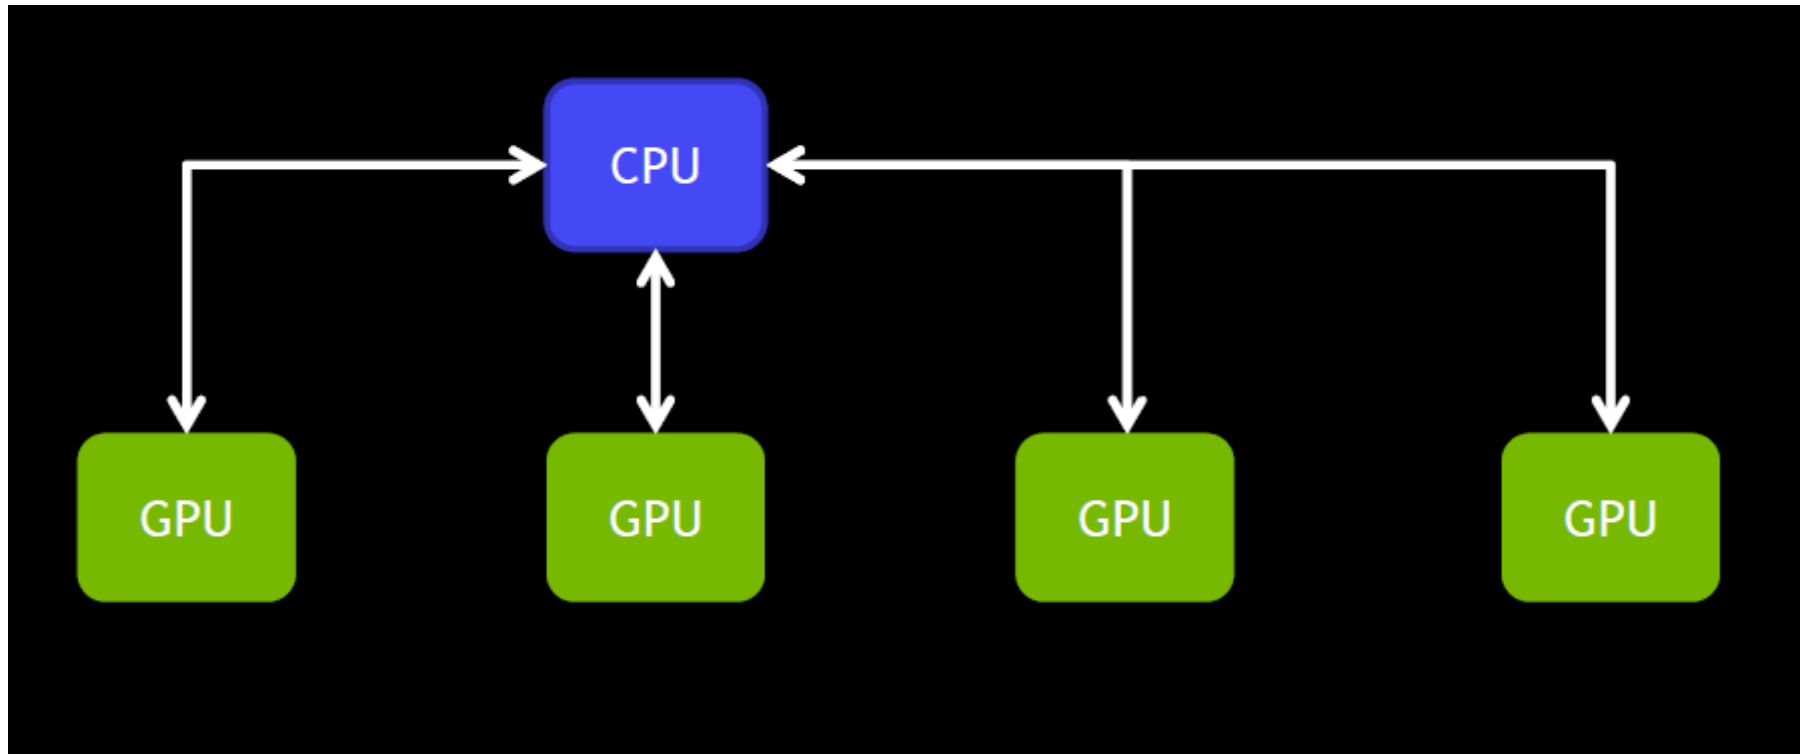
Multi-GPU Programming

Prof. Dr. Alptekin Temizel
Graduate School of Informatics, METU


Single-GPU Case



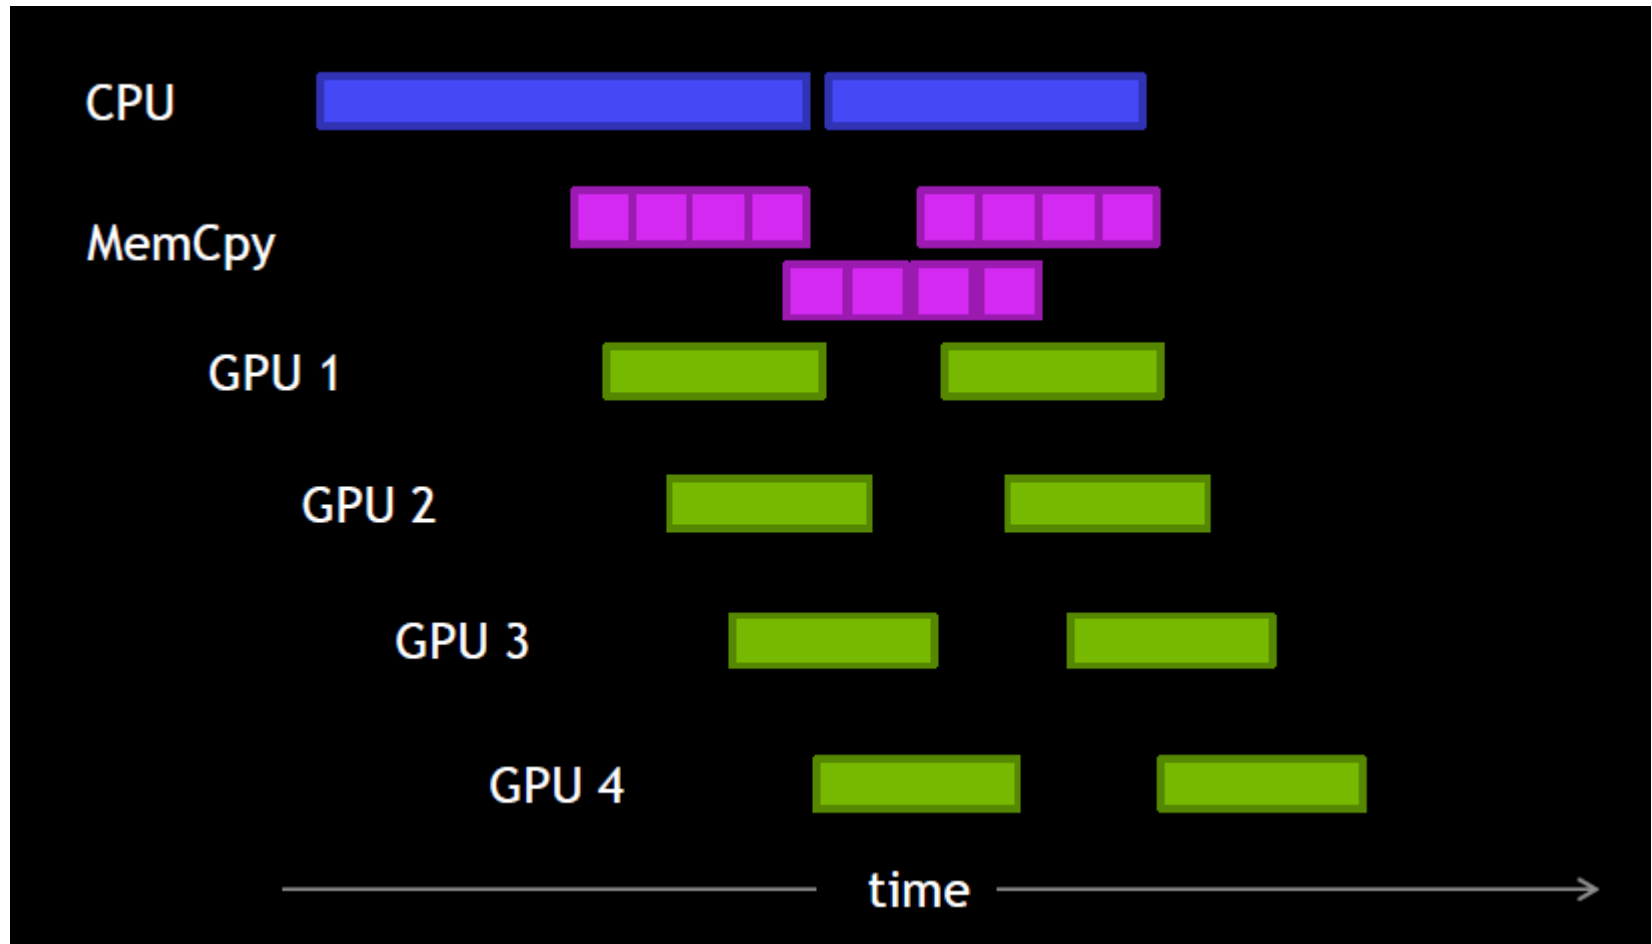
Multi-GPU Case



NVIDIA K10 - Multi-GPU on a Board



Multi-GPU Case



Managing Multiple GPUs

cudaSetDevice() - sets the current GPU

```
cudaSetDevice( 0 );
```

```
kernel<<<...>>>(...);
```

```
cudaMemcpyAsync(...);
```

```
cudaSetDevice( 1 );
```


```
kernel<<<...>>>(...);
```



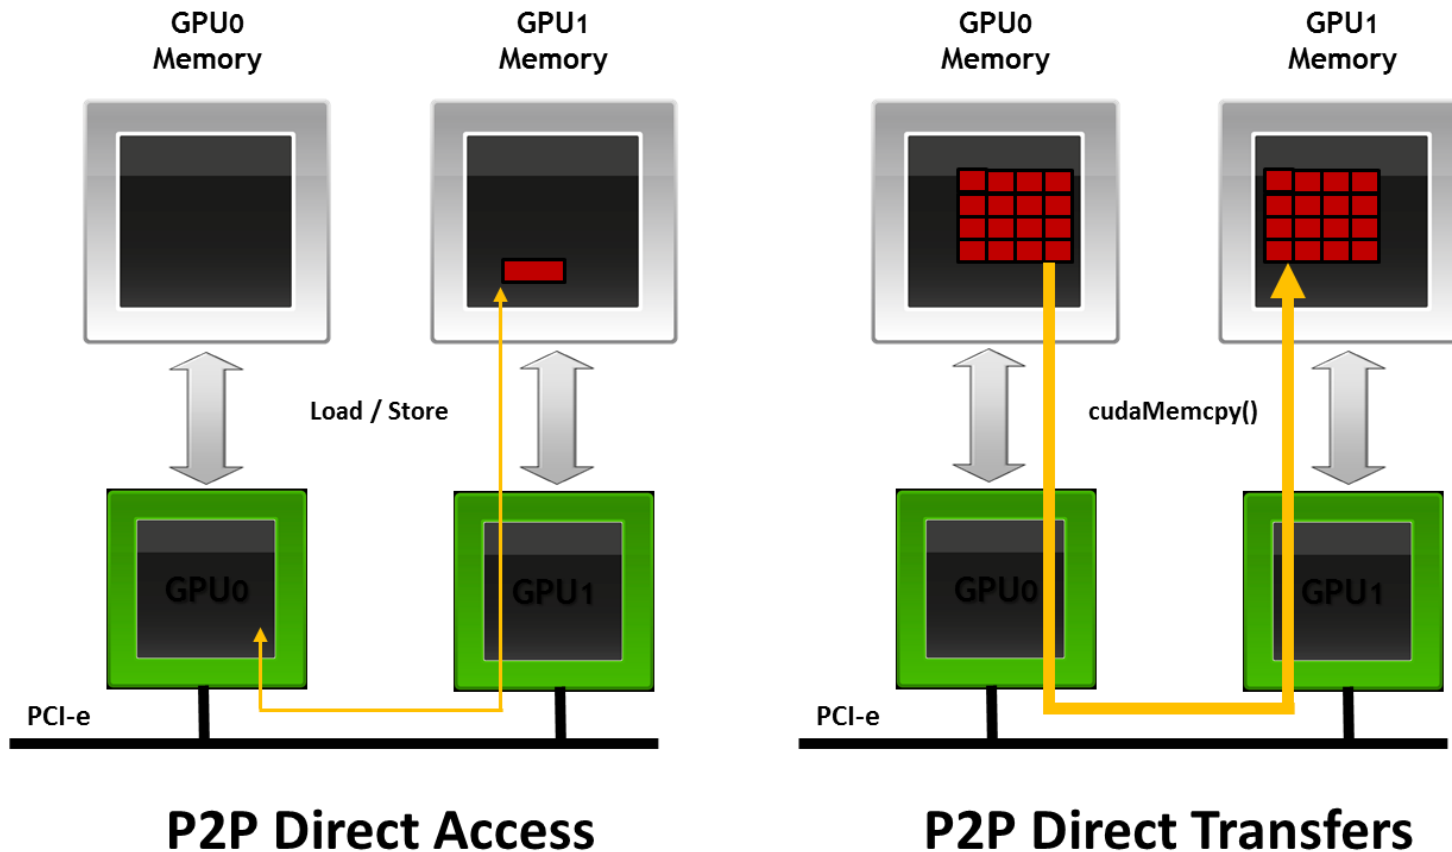
Peer-to-Peer Memcopy

```
cudaMemcpyPeerAsync(  
void* dst_addr, int dst_dev,  
void* src_addr, int src_dev,  
size_t num_bytes,  
cudaStream_t stream )
```

If you analyze with the profiler, you'll see two
separate copies:
D2H and then H2D



GPUDirect Peer-to-Peer (P2P) Communication Between GPUs on the Same PCIe Bus




GPU Direct

Without GPUDirect

Same data copied three times:


1. GPU writes to pinned systemmem1
2. CPU copies from systemmem1 to systemmem2
3. InfiniBand driver copies from systemmem2



With GPUDirect

Data only copied twice

Sharing pinned system memory makes system-to-system copy unnecessary



Enabling GPU Direct

❑ `cudaDeviceEnablePeerAccess(peer_device, 0)`

- Enables current GPU to access addresses on *peer_device* GPU

❑ `cudaDeviceCanAccessPeer(&accessible, dev_X, dev_Y)`

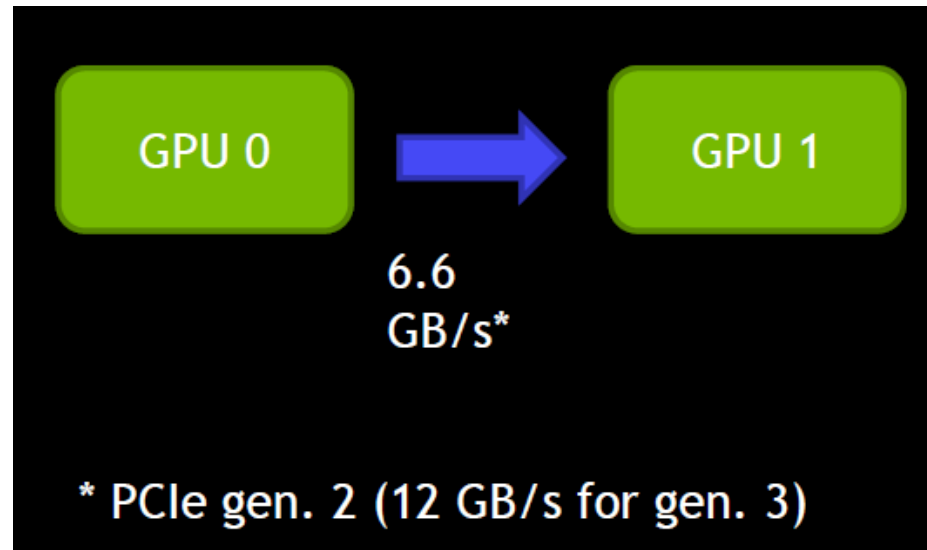
- Checks whether dev_X can access memory of dev_Y

*Need to have at least FermiCUDA 4.0 and same type devices (not possible between Fermi to Kepler etc.)



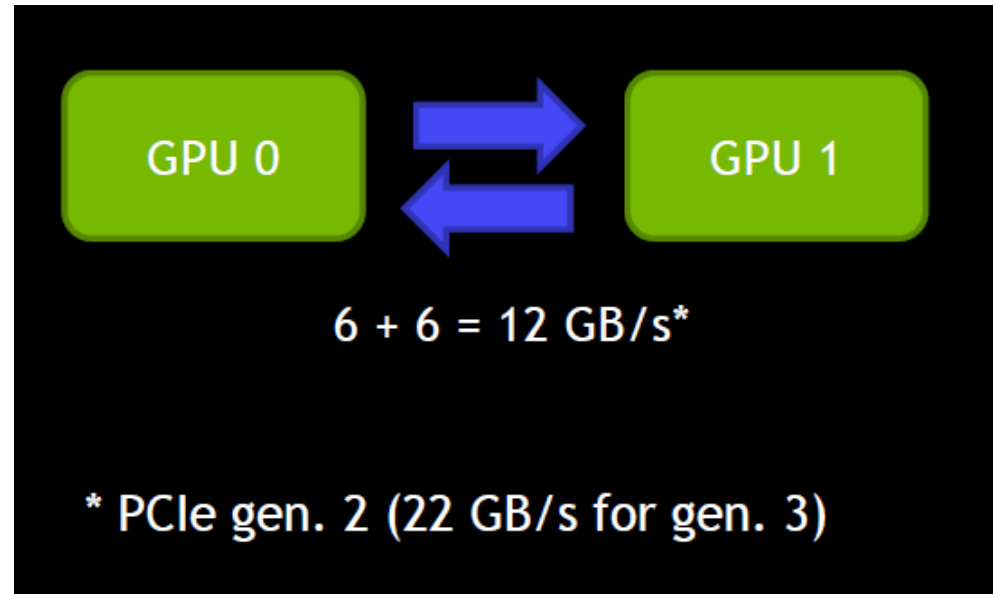
Peer-to-Peer Memcopy with GPU Direct

```
cudaMemcpyPeerAsync(  
void* dst_addr, int dst_dev,  
void* src_addr, int src_dev,  
size_t num_bytes,  
cudaStream_t stream )
```



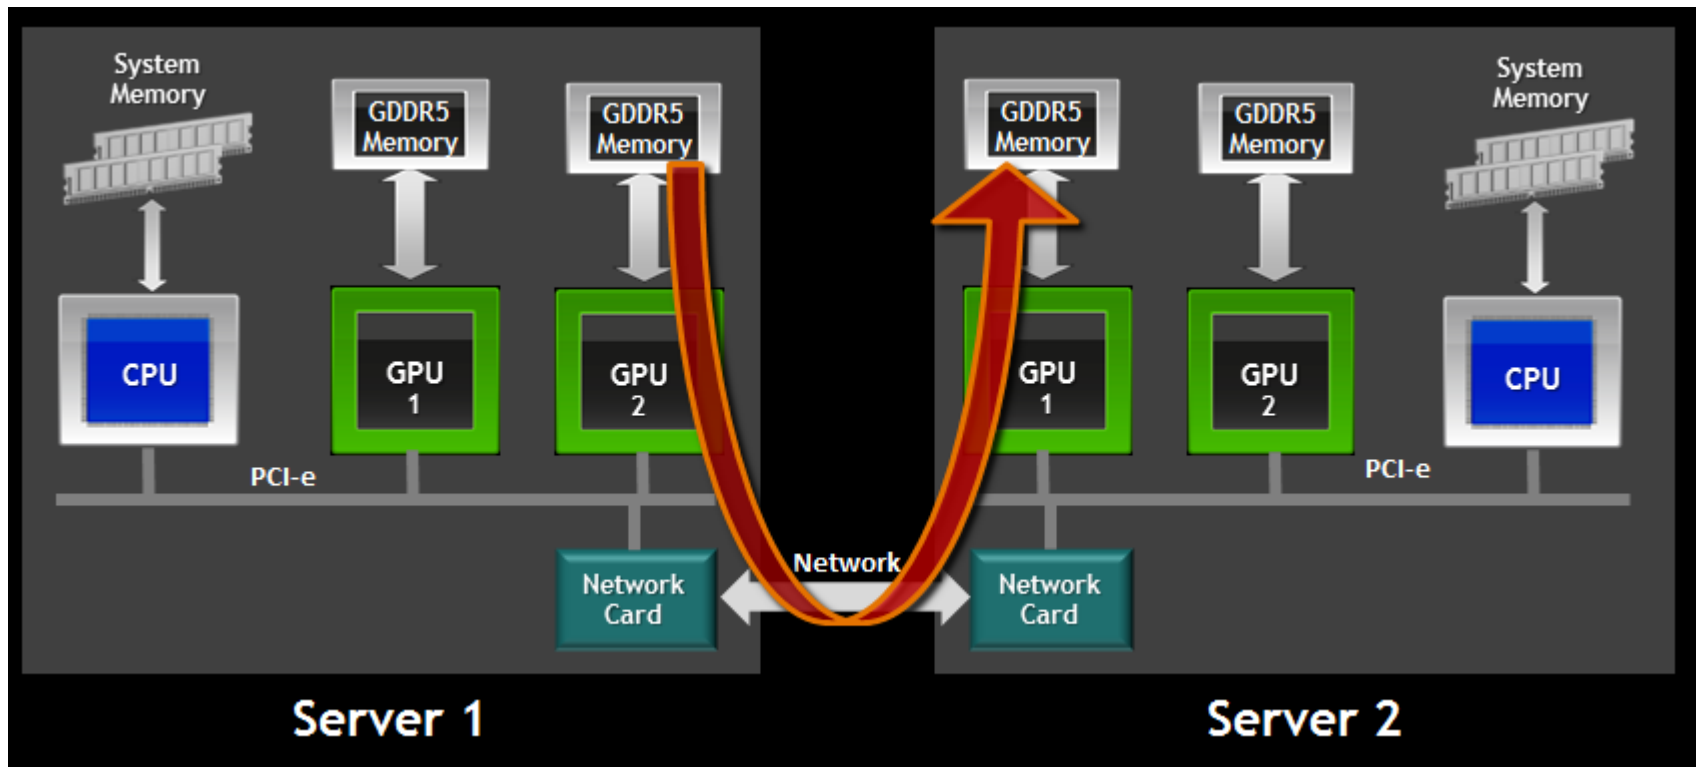
Peer-to-Peer Memcopy with GPU Direct

```
cudaMemcpyPeerAsync(  
void* dst_addr, int dst_dev,  
void* src_addr, int src_dev,  
size_t num_bytes,  
cudaStream_t stream )
```



Peer-to-Peer Memcopy with GPU Direct

- GPU Direct RDMA enables third party PCI Express devices to directly access GPU bypassing CPU host memory altogether.
 - Results in significantly improved MPI SendRecv efficiency between GPUs and other nodes)



Peer-to-Peer Memcopy with GPU Direct

- GPU Direct RDMA enables third party PCI Express devices to directly access GPU bypassing CPU host memory altogether.
 - Results in significantly improved MPI SendRecv efficiency between GPUs and other nodes)
- GPUDirect for Video enables optimized pipeline for frame-based devices such as frame grabbers, video switchers, HD-SDI capture, and CameraLink devices.



Unified Virtual Addressing (UVA)

- ❑ Driver/GPU can determine from an address where data resides (whether in CPU or GPU and on which GPU as well)
 - Requires:
 - 64-bit Linux or 64-bit Windows with TCC (Tesla Compute Cluster) driver
 - Fermi or later architecture GPUs (compute capability 2.0 or higher)
 - CUDA 4.0 or later



Peer-to-peer memcopy


❑ `cudaMemcpyAsync(void* dst_addr, void* src_addr, size_t num_bytes, cudaStream_t stream, cudaMemcpyDefault)`

➤ Just give dst and src addresses and it'll decide on to copy from where to where



Reminder – Overlapping kernel and memcopy

```
cudaStream_t stream1, stream2;  
cudaStreamCreate(&stream1);  
cudaStreamCreate(&stream2);  
cudaHostAlloc(&src, size, 0);  
...  
cudaMemcpyAsync( dst, src, size, dir, stream1 );  
kernel<<<grid, block, 0, stream1>>>(...);  
cudaMemcpyAsync( dst, src, size, dir, stream2 );  
kernel<<<grid, block, 0, stream2>>>(...);
```



Reminder – Creating CUDA Events

- ❑ `cudaEventCreate(cudaEvent_t *event)`
 - Creates an event object
- ❑ `cudaEventDestroy(cudaEvent_t event)`
 - Destroys the specified event
- ❑ `cudaEventRecord(cudaEvent_t event, cudaStream_t stream=0)`
 - Records an event.
 - If stream is non-zero, the event is recorded after all preceding operations in the stream have been completed;
 - Otherwise, it is recorded after all preceding operations in the CUDA context have been completed.
 - `cudaEventQuery()` or `cudaEventSynchronize()` must be used to determine when the event has actually been recorded.




Reminder – Using CUDA Events

- ❑ `cudaEventSynchronize(cudaEvent_t event)`
 - Wait until the completion of all device work preceding the event
- ❑ `cudaStreamWaitEvent(cudaStream_t stream, cudaEvent_t event)`
 - Makes all future work submitted to the stream wait until event reports completion before beginning execution.
 - This synchronization will be performed efficiently on the device.



Reminder – Using CUDA Events

```
cudaEvent_t ev;  
cudaEventCreate(&ev);  
...  
cudaMemcpyAsync( dst, src, size, dir, stream1 );  
cudaMemcpyAsync( dst2, src2, size, dir, stream2 );  
cudaEventRecord( ev, stream2 );  
cudaStreamWaitEvent(stream1, ev);  
kernel<<<grid, block, 0, stream1>>>(...);  
kernel<<<grid, block, 0, stream2>>>(...);
```



Multi-GPU Streams and Events

- ❑ CUDA streams and events are *per device* (GPU)
 - Each device has its own *default* stream (stream 0)
- ❑ **Kernels:** can be launched to a stream only if the stream's GPU is current
- ❑ **Memcopies:** can be issued to any stream
- ❑ **Events:** can be recorded only to a stream if the stream's GPU is current
- ❑ It is OK to query or synchronize with any event/stream



Multi-GPU Streams and Events

```
cudaStream_t streamA, streamB;  
cudaEvent_t eventA, eventB;  
cudaSetDevice( 0 );  
cudaStreamCreate( &streamA ); // streamA and eventA belong to device-0  
cudaEventCreate( &eventA );  
cudaSetDevice( 1 );  
cudaStreamCreate( &streamB ); // streamB and eventB belong to device-1  
cudaEventCreate( &eventB );  
kernel<<<..., streamB>>>(...);  
cudaEventRecord( eventB, streamA );  
cudaEventSynchronize( eventB );
```

NOT OK:

- device 1 is current
- streamA belongs to device 0



Multi-GPU Streams and Events

```
cudaStream_t streamA, streamB;  
cudaEvent_t eventA, eventB;  
cudaSetDevice( 0 );  
cudaStreamCreate( &streamA ); // streamA and eventA belong to device-0  
cudaEventCreate( &eventA );  
cudaSetDevice( 1 );  
cudaStreamCreate( &streamB ); // streamB and eventB belong to device-1  
cudaEventCreate( &eventB );  
kernel<<<..., streamA>>>( ... );  
cudaEventRecord( eventB, streamB );  
cudaEventSynchronize( eventB );
```

NOT OK:

- device 1 is current
- streamA belongs to device 0



Multi-GPU Streams and Events

```
cudaStream_t streamA, streamB;  
cudaEvent_t eventA, eventB;  
cudaSetDevice( 0 );  
cudaStreamCreate( &streamA ); // streamA and eventA belong to device-0  
cudaEventCreate( &eventA );
```

```
cudaSetDevice( 1 );  
cudaStreamCreate( &streamB ); // streamB and eventB belong to device-1  
cudaEventCreate( &eventB );  
kernel<<<..., streamB>>>(...);  
cudaEventRecord( eventB, streamB );
```

```
cudaSetDevice( 0 );  
cudaEventSynchronize( eventB );  
kernel<<<..., streamA>>>(...);
```

- Synchronizing/querying events/streams of other devices is allowed




Multi-GPU Streams and Events

```
cudaSetDevice( 0 );  
cudaMalloc( &d_A, num_bytes );  
int accessible = 0;  
cudaDeviceCanAccessPeer( &accessible, 1, 0 );  
if( accessible )  
{  
    cudaSetDevice( 1 );  
    cudaDeviceEnablePeerAccess( 0, 0 );  
    kernel<<<...>>>( d_A );  
}
```


- ❑ Kernel executes on Device 1
- ❑ It will access (via PCIe) memory on Device 0




Subdividing the Problem




Subdividing the Problem



Step 1
Compute Halo Regions



Step 2
Compute Internal Regions



Step 3
Exchange Halo Regions



Subdividing the Problem

Code Pattern

```
for( int istep=0; istep<nsteps; istep++)  
{  
    for( int i=0; i<num_gpus; i++ )  
    {  
        cudaSetDevice( gpu[i] );  
        kernel_halo<<<..., stream_compute[i]>>>( ... );  
  
        kernel_int<<<..., stream_compute[i]>>>( ... );  
    }  
    for( int i=0; i<num_gpus-1; i++ ) {  
  
        cudaMemcpyPeerAsync( ..., stream_compute[i] );  
    }  
    for( int i=1; i<num_gpus; i++ )  
        cudaMemcpyPeerAsync( ..., stream_compute[i] );  
}
```

Compute halos

Compute internal

Exchange halos

stream_compute

halo

int

halo

int



Subdividing the Problem

Code Pattern

```
for( int istep=0; istep<nsteps; istep++)
```

```
{
```

```
  for( int i=0; i<num_gpus; i++ )
```

```
  {
```

```
    cudaSetDevice( gpu[i] );
```

```
    kernel_halo<<<..., stream_compute[i]>>>( ... );
```

```
    cudaEventRecord(event_i[i], stream_compute[i] );
```

```
    kernel_int<<<..., stream_compute[i]>>>( ... );
```

```
  }
```

```
  for( int i=0; i<num_gpus-1; i++ ) {
```

```
    cudaStreamWaitEvent(stream_copy[i], event_i[i]);
```

```
    cudaMemcpyPeerAsync( ..., stream_copy[i] );
```

```
  }
```

```
  for( int i=1; i<num_gpus; i++ )
```

```
    cudaMemcpyPeerAsync( ..., stream_copy[i] );
```

```
}
```

Compute halos

Compute internal

Exchange halos

stream_compute

stream_copy

halo

int

halo

int

halo

int



Subdividing the Problem

Code Pattern

```
for( int istep=0; istep<nsteps; istep++)  
{
```

```
  for( int i=0; i<num_gpus; i++ )
```

```
  {
```

```
    cudaSetDevice( gpu[i] );
```

```
    kernel_halo<<<..., stream_compute[i]>>>( ... );
```

```
    cudaEventRecord(event_i[i], stream_compute[i] );
```

```
    kernel_int<<<..., stream_compute[i]>>>( ... );
```

```
  }
```

```
  for( int i=0; i<num_gpus-1; i++ ) {
```

```
    cudaStreamWaitEvent(stream_copy[i], event_i[i]);
```

```
    cudaMemcpyPeerAsync( ..., stream_copy[i] );
```

```
  }
```

```
  for( int i=1; i<num_gpus; i++ )
```

```
    cudaMemcpyPeerAsync( ..., stream_copy[i] );
```

```
  for( int i=0; i<num_gpus; i++ )
```

```
  {
```

```
    cudaSetDevice( gpu[i] );
```

```
    cudaDeviceSynchronize();
```

```
    // swap input/output pointers
```

```
  }
```

```
}
```

Compute halos

Compute internal


Exchange halos

Synchronize



MPI

- ❑ MPI (Message Passing Interface) is a standardized and portable API for communicating data via messages (both point-to-point & collective) between distributed processes.
- ❑ MPI is frequently used in HPC to build applications that can scale on multi-node computer clusters.
- ❑ In most MPI implementations, library routines are directly callable from C, C++, and Fortran.



Combining MPI with CUDA

- ☐ To solve problems with a data size that is too large to fit into the memory of a single GPU.
- ☐ To solve problems that would require unreasonably long compute time on a single node.
- ☐ To accelerate an existing MPI application with GPUs.
- ☐ To enable a single-node multi-GPU application to scale across multiple nodes.



CUDA Aware MPI

- ❑ Regular MPI implementations pass pointers to host memory, staging GPU buffers through host memory using `cudaMemcpy`.
- ❑ With CUDA-aware MPI, the MPI library can send and receive GPU buffers directly, without having to first stage them in host memory.
- ❑ Implementation of CUDA-aware MPI was simplified by Unified Virtual Addressing (UVA) in CUDA 4.0 – which enables a single address space for all CPU and GPU memory.



CUDA Aware MPI


- ❑ MVAPICH2 is an open source MPI implementation
 - Simplifies the task of porting MPI applications to run on clusters with NVIDIA GPUs by supporting standard MPI calls from GPU device memory.

- ❑ IBM® Platform™ MPI is a high performance, production-quality MPI implementation
 - Supports a broad range of industry standard platforms, interconnects and operating systems.

- ❑ The Open MPI Project is an open source MPI-2 implementation that is developed and maintained by a consortium of academic, research, and industry partners.
 - GPUs are supported by version 1.7 and later.



Multiple GPU, Multiple Nodes



```
cudaMemcpyAsync( ..., stream_halo[i] );  
cudaStreamSynchronize( stream_halo[i] );  
MPI_Sendrecv( ... );  
cudaMemcpyAsync( ..., stream_halo[i] );
```



NVLink

- NVLink is NVIDIA's new high-speed data interconnect. NVLink can be used to significantly increase performance for both GPU-to-GPU communication and for GPU access to system memory. GP100 supports up to four NVLink connections with each connection carrying up to 40 GB/s of bi-directional bandwidth.
- NVLink is available between GPU to GPU transfers
 - x86 devices: CPU to GPU transfers are through PCIe
 - Power8/Power9: CPU to GPU transfers are also through NVLink

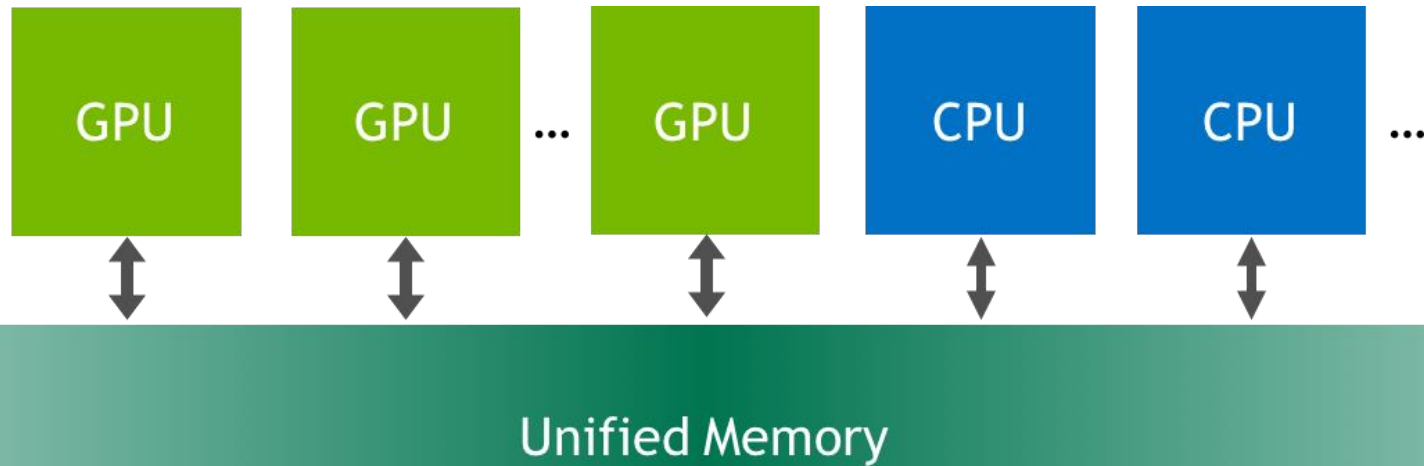


NVLink

- NVLink operates transparently within the existing CUDA model. Transfers between NVLink-connected endpoints are automatically routed through NVLink, rather than PCIe.
- The `cudaDeviceEnablePeerAccess()` API call remains necessary to enable direct transfers (over either PCIe or NVLink) between GPUs. The `cudaDeviceCanAccessPeer()` can be used to determine if peer access is possible between any pair of GPUs.



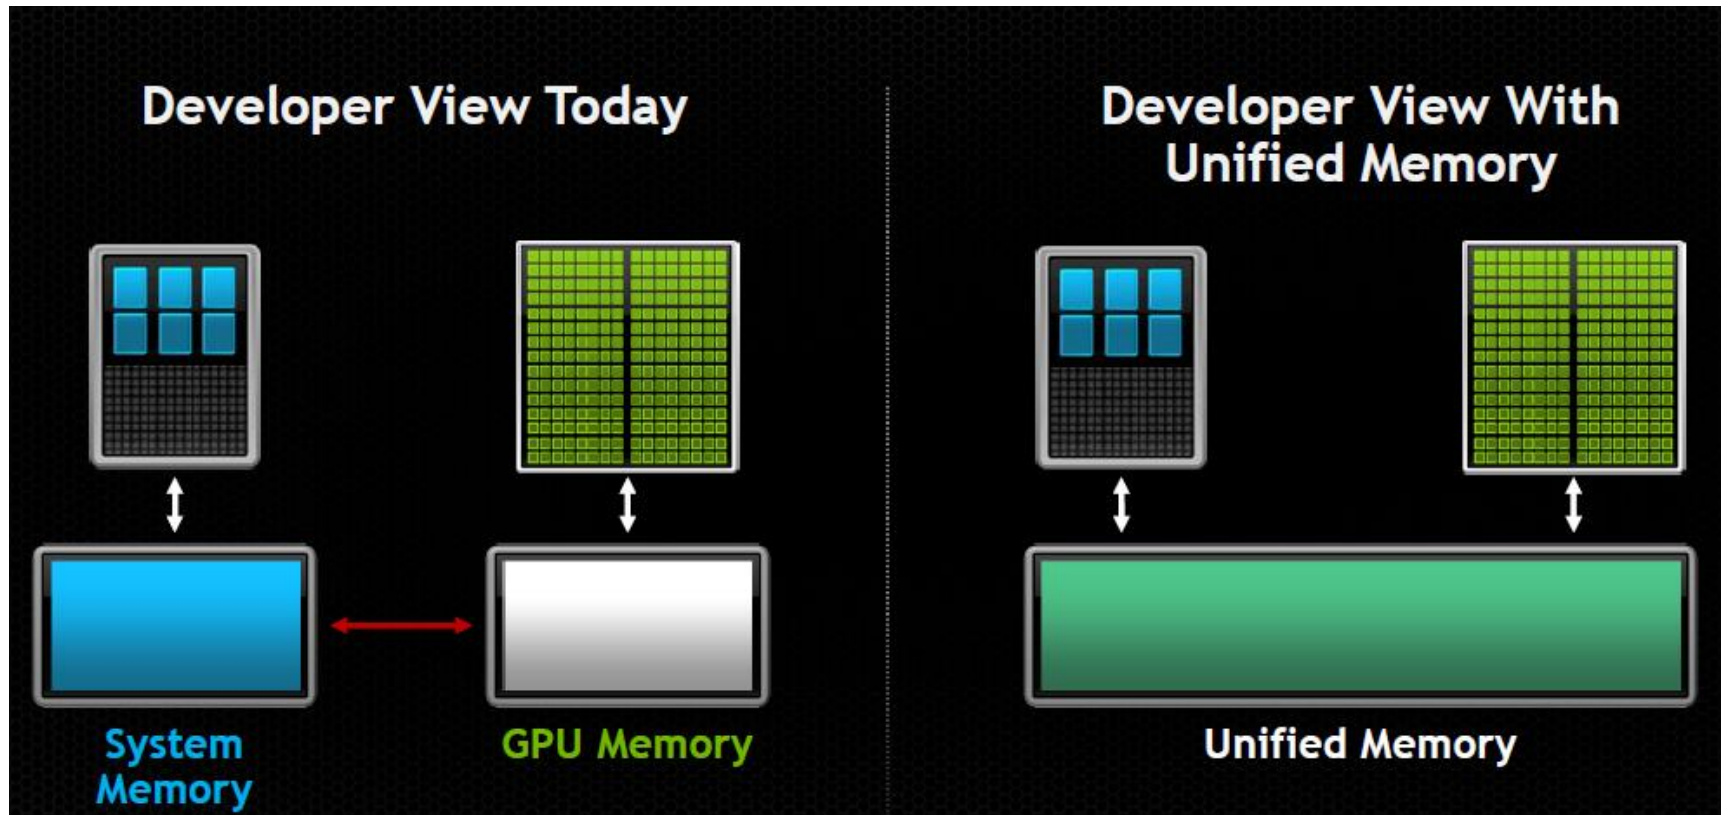
Unified Memory



Unified Memory is a single memory address space accessible from any processor in a system



Unified Memory



Unified Memory

CPU Version

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    data = (char *)malloc(N);  
    fread(data, 1, N, fp);  
    qsort(data, N, 1, compare);  
    use_data(data);  
    free(data);  
}
```

GPU Unified Memory Version

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    cudaMallocManaged(&data, N);  
    fread(data, 1, N, fp);  
    qsort<<<...>>>(data, N, 1, compare);  
    cudaDeviceSynchronize();  
    use_data(data);  
    cudaFree(data);  
}
```



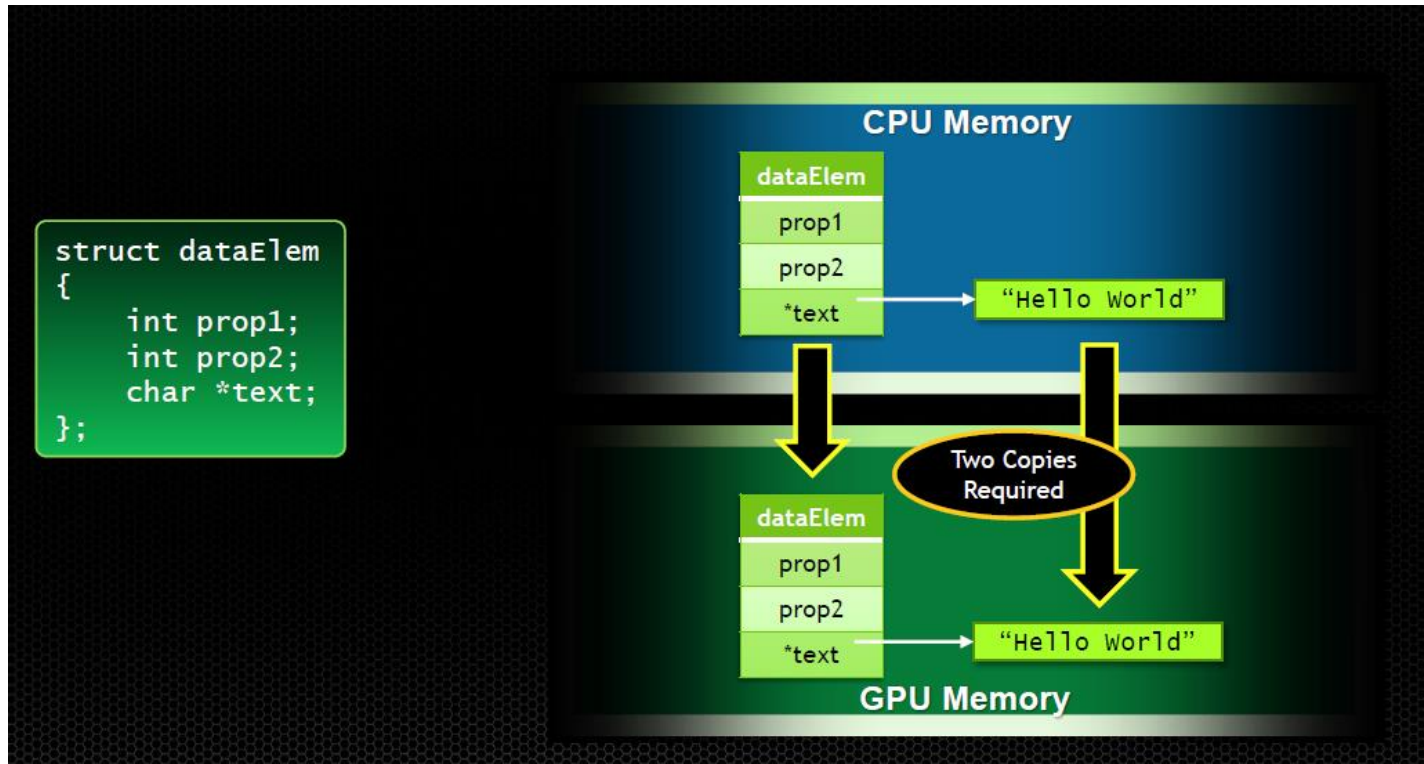
Unified Memory

- ❑ Single pointer to data, accessible anywhere
- ❑ Guarantee global coherency
- ❑ Still allows *cudaMemcpyAsync()* hand tuning
- ❑ Eliminates the need for deep copying when using structs



Unified Memory

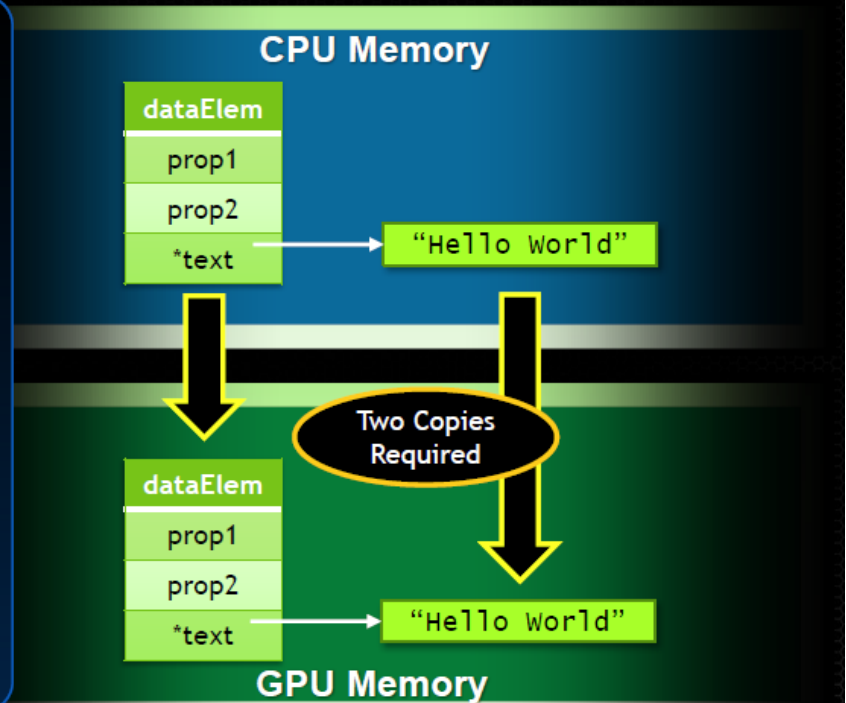
- ❑ Eliminates the need for deep copying when using structs



Unified Memory

- ❑ Eliminates the need for deep copying when using structs

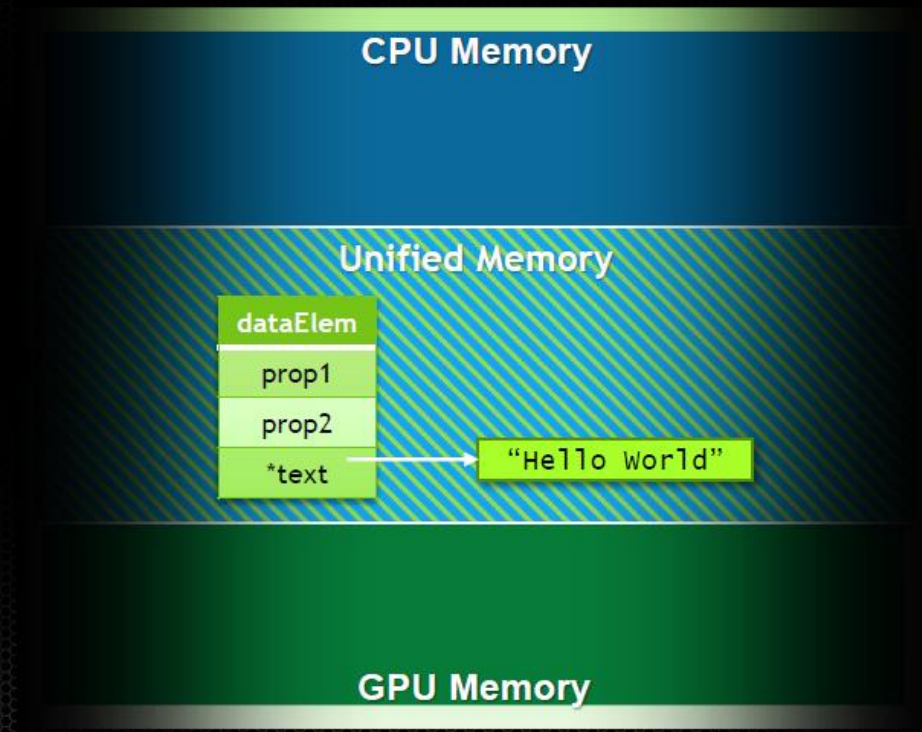
```
void launch(dataElem *elem) {  
    dataElem *g_elem;  
    char *g_text;  
  
    int textlen = strlen(elem->text);  
  
    // Allocate storage for struct and text  
    cudaMalloc(&g_elem, sizeof(dataElem));  
    cudaMalloc(&g_text, textlen);  
  
    // Copy up each piece separately, including  
    // new "text" pointer value  
    cudaMemcpy(g_elem, elem, sizeof(dataElem));  
    cudaMemcpy(g_text, elem->text, textlen);  
    cudaMemcpy(&(g_elem->text), &g_text,  
               sizeof(g_text));  
  
    // Finally we can launch our kernel, but  
    // CPU & GPU use different copies of "elem"  
    kernel<<< ... >>>(g_elem);  
}
```



Unified Memory

- ❑ Eliminates the need for deep copying when using structs

```
void launch(dataElem *elem) {  
    kernel<<< ... >>>(elem);  
}
```



Unified Memory

Example: GPU & CPU Shared Linked Lists

- Can pass list elements between Host & Device
- Can insert and delete elements from Host or Device*
- Single list - no complex synchronization

*Program must still ensure no race conditions.
Data is coherent between CPU & GPU
at kernel launch & sync only

